

# CUDA C++ Exercise: Basic Linear Algebra Kernels: GEMM Optimization Strategies

Dmitry Lyakh

Scientific Computing

Oak Ridge Leadership Computing Facility

Oak Ridge National Laboratory

ORNL is managed by UT-Battelle, LLC for the US Department of Energy

# Installing CUDA Basic Linear Algebra (BLA) Library

- Log in to BlueWaters @ NCSA:  
**ssh <user\_id>@bwbay.ncsa.illinois.edu**
- Clone the BLA exercise repository (either way):  
**git clone [https://github.com/DmitryLyakh/CUDA\\_Tutorial.git](https://github.com/DmitryLyakh/CUDA_Tutorial.git)**  
or  
**git clone /projects/training/bayr/CUDA\_Tutorial**
- If already cloned, **cd CUDA\_Tutorial** and **git pull**
- Adjust Cray environment modules:  
**module swap PrgEnv-cray PrgEnv-gnu**  
**module swap gcc/8.2.0**  
**module load cudatoolkit**
- Copy Makefile: **/projects/training/bayr/CUDA\_Tutorial/Makefile** into your CUDA\_Tutorial directory or make sure CUDA\_HOST, CUDA\_INC, CUDA\_LIB match in the Makefile
- Run **make**: Builds binary **bla\_test.x**
- Running **bla\_test.x** on BlueWaters:
  - Open interactive session **once** (for one hour): **qsub -I -l nodes=1:ppn=16:xk -l walltime=01:00:00**
  - Adjust Cray modules again as described above (**once** per interactive session)
  - **cd CUDA\_Tutorial** (enter your CUDA\_Tutorial path)
  - **aprun -n1 -N1 -d16 ./bla\_test.x**

# CUDA BLA Library Concepts: Matrix

- In file matrix.hpp: **class Matrix<T>, T = {float, double}**
- Matrix constructor: **Matrix(nrows, ncols)**: No storage yet!
- Matrix storage: **Matrix.allocateBody(int device)**:  
CPU Host: device = -1  
NVIDIA GPU: device = 0,1,2,... (only one GPU on BlueWaters)  
May simultaneously reside on Host and GPU: Needs sync (below)!
- Set to zero on given device: **Matrix.zeroBody(int device)**
- Set to some random value on Host: **Matrix.setBodyHost()**
- Synchronize value on multiple devices:  
**Matrix.syncBody(int device, int source\_device)**

# CUDA BLA Library Concepts: Matrix Operations

- Compute sum of the squares of all elements (on given device):  
**double Matrix.computeNorm(int device)**
- Add one matrix to another matrix (on given device):  
**Matrix.add(Matrix & Amat, T alpha, int device)**
- Multiply two matrices and add the result to another matrix:  
**Matrix.multiplyAdd(bool left\_transp, bool right\_transp,  
Matrix & Amat, Matrix & Bmat, int device)**
- Default Matrix.multiplyAdd GPU implementation expects:  
left\_transp = **false**, right\_transp = **false**
- **Your exercise is to implement GPU kernels for all transposition cases: FalseTrue, TrueFalse, TrueTrue**

# CUDA BLA Library Implementation Benchmark

- Our test driver code: main.cpp: Function **use\_bla()**
- Creates matrices **A(m,k)**, **B(k,n)**, **C(m,n)** with some **m**, **n**, **k**
- Computes the total **flop count** for matrix multiplication:  
**Flop = 2\*m\*n\*k**, where factor of 2 is (multiply + add) = 2 Flop
- Executes matrix multiplication/accumulation (**C.multiplyAdd**):  
**C(m,n) += A(m,k) \* B(k,n)**
- Function **bla::reset\_gemm\_algorithm(int)** chooses between:
  - 7: Highly optimized cuBLAS GEMM implementation
  - 2: Shared memory + registers based BLA GEMM (bla\_lib.cu: **gpu\_gemm\_sh\_reg\_nn**)
  - 1: Shared memory based BLA GEMM (bla\_lib.cu: **gpu\_gemm\_sh\_nn**)
  - 0: Naïve BLA GEMM implementation (bla\_lib.cu: **gpu\_gemm\_nn**)

# CUDA BLA Library Implementation Benchmark

Testing your BLA GPU kernel implementation (main.cpp: use\_bla() function):

```
for(int repeat = 0; repeat < 2; ++repeat){ //repeat experiment twice
    C.zeroBody(0); //set matrix C body to zero on GPU#0
    bla::reset_gemm_algorithm(0); //choose your algorithm: {0,1,2,7}
    std::cout << "Performing matrix multiplication C+=A*B with BLA GEMM brute-force ... ";
    double tms = bla::time_sys_sec(); //timer start
    C.multiplyAdd(false,false,A,B,0); //default case {false,false}: You goal is {false,true}, {true,false}, {true,true}
    double tmf = bla::time_sys_sec(); //timer stop
    std::cout << "Done: Time = " << tmf-tms << " s: Gflop/s = " << flops/(tmf-tms)/1e9 << std::endl;
    //Check correctness on GPU#0:
    C.add(D,1.0f,0); //adding the correct result with a minus sign (matrix D) should give you zero matrix
    auto norm_diff = C.computeNorm(0); //check its norm
    std::cout << "Norm of the matrix C deviation from correct = " << norm_diff << std::endl;
    if(std::abs(norm_diff) > 1e-7){ //report if norm is not zero enough
        std::cout << "#FATAL: Matrix C is incorrect, fix your GPU kernel implementation!" << std::endl;
        std::exit(1);
    }
}
```

This benchmark is run for all available BLA GEMM algorithms: 0, 1, 2, 7 for the {false,false} case. Your goal is to implement and run other cases: {false,true}, {true,false}, {true,true}!

# CUDA BLA Library: GEMM cases

- **Matrix A(m,n)** employs column-wise storage (standard BLAS):
  - $A(0,0), A(1,0), A(2,0), \dots, A(m-1,0), A(0,1), A(1,1), A(2,1), \dots, A(m-1,n-1)$
  - **Linear offset** of element (j,k) in storage is  $L(j,k) = (j + k*m)$
  - **m** is the leading dimension extent in this case
- Matrix multiplication {**false,false**} case (implemented):
  - $C(m,n) += A(m,k) * B(k,n)$
- Matrix multiplication {**false,true**} case (your exercise):
  - $C(m,n) += A(m,k) * B(n,k)$
- Matrix multiplication {**true,false**} case (your exercise):
  - $C(m,n) += A(k,m) * B(k,n)$
- Matrix multiplication {**true,true**} case (your exercise):
  - $C(m,n) += A(k,m) * B(n,k)$



# CUDA BLA Library: GEMM algorithms

- You will work inside **bla\_lib.cu** source file directly with CUDA GEMM kernels
- Matrix multiplication **{false,false}** case (implemented):
  - $C(m,n) += A(m,k) * B(k,n)$
  - CUDA kernels: **gpu\_gemm\_nn, gpu\_gemm\_sh\_nn, gpu\_gemm\_sh\_reg\_nn**
- Matrix multiplication **{false,true}** case (your exercise):
  - $C(m,n) += A(m,k) * B(n,k)$
  - CUDA kernels: **gpu\_gemm\_nt, gpu\_gemm\_sh\_nt, gpu\_gemm\_sh\_reg\_nt**
- Matrix multiplication **{true,false}** case (your exercise):
  - $C(m,n) += A(k,m) * B(k,n)$
  - CUDA kernels: **gpu\_gemm\_tn, gpu\_gemm\_sh\_tn, gpu\_gemm\_sh\_reg\_tn**
- Matrix multiplication **{true,true}** case (your exercise):
  - $C(m,n) += A(k,m) * B(n,k)$
  - CUDA kernels: **gpu\_gemm\_tt, gpu\_gemm\_sh\_tt, gpu\_gemm\_sh\_reg\_tt**



# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```
template <typename T>
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
{
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;

    size_t n_pos = ty;
    while(n_pos < n){

        size_t m_pos = tx;
        while(m_pos < m){

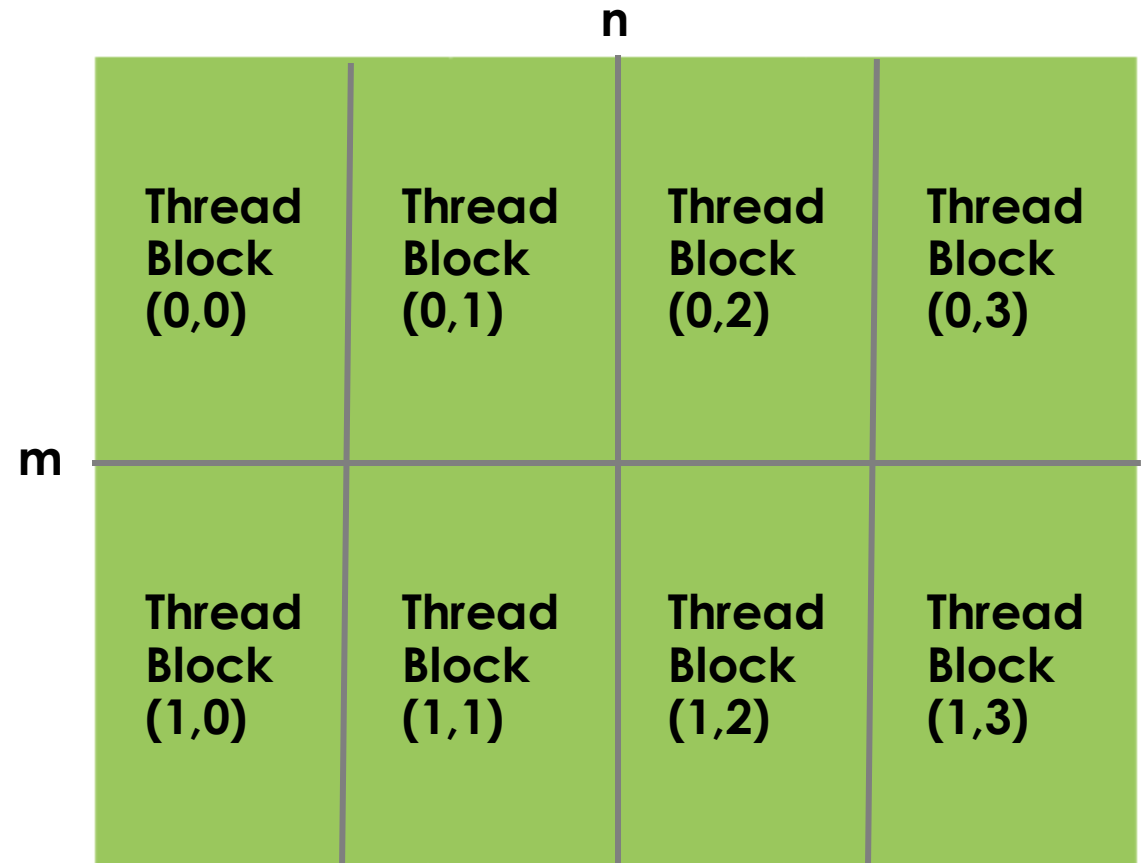
            T tmp = static_cast<T>(0.0);
            for(size_t k_pos = 0; k_pos < k; ++k_pos){
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];
            }
            dest[n_pos*m + m_pos] += tmp;

            m_pos += blockDim.x*blockDim.x;
        }

        n_pos += blockDim.y*blockDim.y;
    }
    return;
}
```

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$



Matrix  $C(m, n)$

# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```
template <typename T> float or double
```

```
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
```

```
{  
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;  
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;
```

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

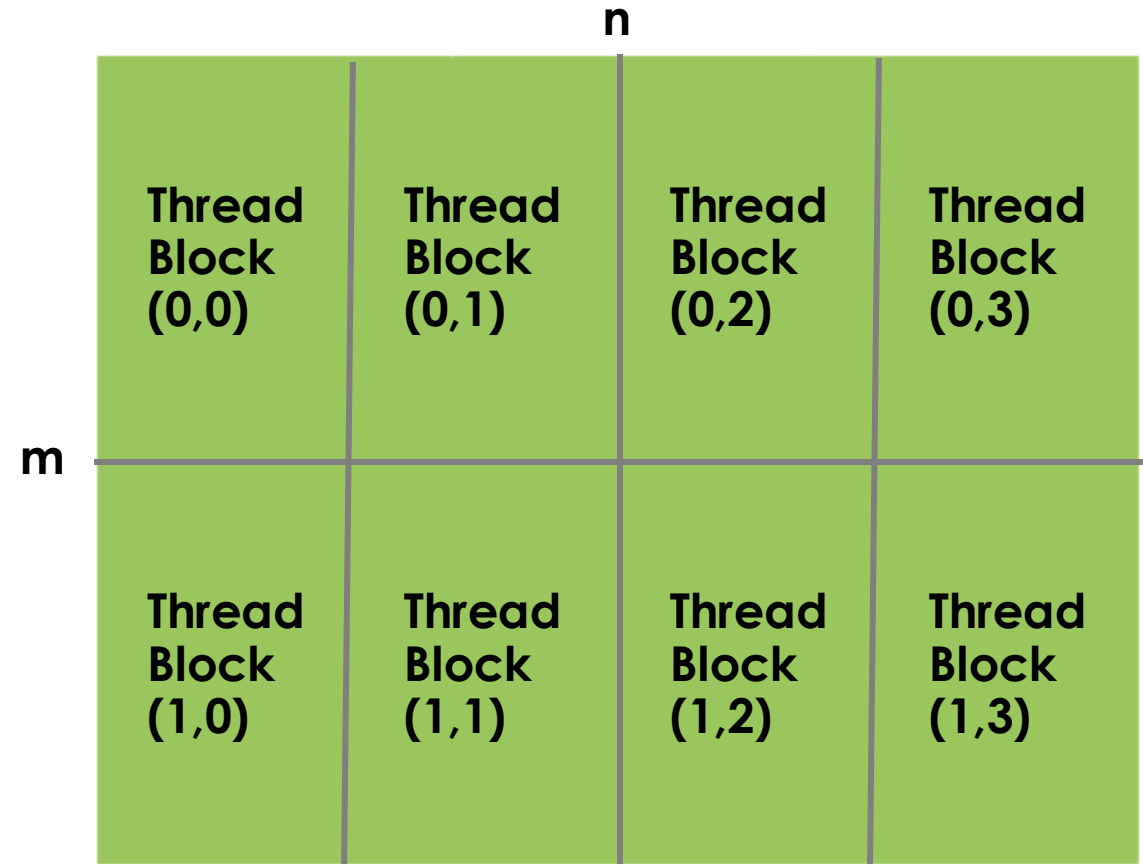
```
    size_t n_pos = ty;  
    while(n_pos < n){
```

```
        size_t m_pos = tx;  
        while(m_pos < m){
```

```
            T tmp = static_cast<T>(0.0);  
            for(size_t k_pos = 0; k_pos < k; ++k_pos){  
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];  
            }  
            dest[n_pos*m + m_pos] += tmp;
```

```
            m_pos += blockDim.x*blockDim.x;  
        }
```

```
        n_pos += blockDim.y*blockDim.y;  
    }  
    return;  
}
```



Matrix C(m,n)

# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```
template <typename T>
```

No pointer aliasing

```
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
```

```
{  
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;  
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;
```

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

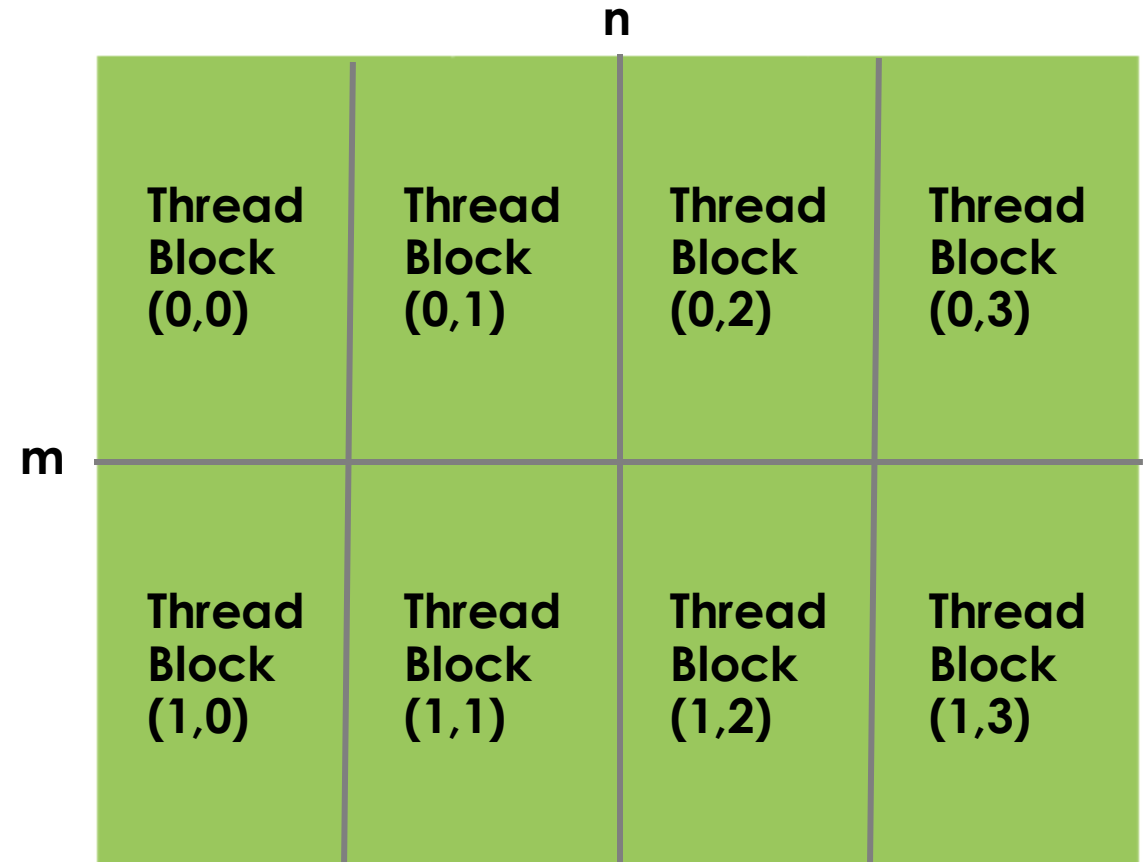
```
    size_t n_pos = ty;  
    while(n_pos < n){
```

```
        size_t m_pos = tx;  
        while(m_pos < m){
```

```
            T tmp = static_cast<T>(0.0);  
            for(size_t k_pos = 0; k_pos < k; ++k_pos){  
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];  
            }  
            dest[n_pos*m + m_pos] += tmp;
```

```
            m_pos += blockDim.x*blockDim.x;  
        }
```

```
        n_pos += blockDim.y*blockDim.y;  
    }  
    return;  
}
```



Matrix C(m,n)

# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```
template <typename T>
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
{
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;

    size_t n_pos = ty;
    while(n_pos < n){

        size_t m_pos = tx;
        while(m_pos < m){

            T tmp = static_cast<T>(0.0);
            for(size_t k_pos = 0; k_pos < k; ++k_pos){
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];
            }
            dest[n_pos*m + m_pos] += tmp;

            m_pos += blockDim.x*blockDim.x;
        }

        n_pos += blockDim.y*blockDim.y;
    }
    return;
}
```

Each CUDA thread block computes:  
 $C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

(7,1)

(13,0)

Matrix C(m,n)

# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```
template <typename T>
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
{
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;

    size_t n_pos = ty;
    while(n_pos < n){
        size_t m_pos = tx;
        while(m_pos < m){
            T tmp = static_cast<T>(0.0);
            for(size_t k_pos = 0; k_pos < k; ++k_pos){
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];
            }
            dest[n_pos*m + m_pos] += tmp;

            m_pos += blockDim.x*blockDim.x;
        }

        n_pos += blockDim.y*blockDim.y;
    }
    return;
}
```

**Each CUDA thread block computes:**  
 $C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

**Bounds guards**

# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```
template <typename T>
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
{
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;

    size_t n_pos = ty;
    while(n_pos < n){

        size_t m_pos = tx;
        while(m_pos < m){

            Init accumulator register
            T tmp = static_cast<T>(0.0);
            for(size_t k_pos = 0; k_pos < k; ++k_pos){
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];
            }
            dest[n_pos*m + m_pos] += tmp;

            m_pos += blockDim.x*blockDim.x;
        }

        n_pos += blockDim.y*blockDim.y;
    }
    return;
}
```

Each CUDA thread block computes:  
 $C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```
template <typename T>
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
{
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;

    size_t n_pos = ty;
    while(n_pos < n){

        size_t m_pos = tx;
        while(m_pos < m){

            T tmp = static_cast<T>(0.0); Loop over entire k dim
            for(size_t k_pos = 0; k_pos < k; ++k_pos){
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];
            }
            dest[n_pos*m + m_pos] += tmp;

            m_pos += blockDim.x*blockDim.x;
        }

        n_pos += blockDim.y*blockDim.y;
    }
    return;
}
```

Each CUDA thread block computes:  
 $C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

The diagram illustrates the layout of Matrix C(m,n) as a 2x4 grid of thread blocks. The vertical axis is labeled 'm' and the horizontal axis is labeled 'n'. The thread blocks are labeled as follows:

- Top row: Thread Block (0,0), Thread Block (0,1), Thread Block (0,2), Thread Block (0,3)
- Bottom row: Thread Block (1,0), Thread Block (1,1), Thread Block (1,2), Thread Block (1,3)

Arrows indicate the dimensions of the thread blocks:

- A red arrow labeled (7,1) points to the top row of thread blocks.
- A red arrow labeled (13,0) points to the left column of thread blocks.

The thread block (0,1) is highlighted with a red dot, indicating the current thread block being computed.



# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```

template <typename T>
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
{
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;

    size_t n_pos = ty;
    while(n_pos < n){

        size_t m_pos = tx;
        while(m_pos < m){

            T tmp = static_cast<T>(0.0);
            for(size_t k_pos = 0; k_pos < k; ++k_pos){
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];
            }
            dest[n_pos*m + m_pos] += tmp;

            m_pos += blockDim.x;
        }

        n_pos += blockDim.y;
    }
    return;
}

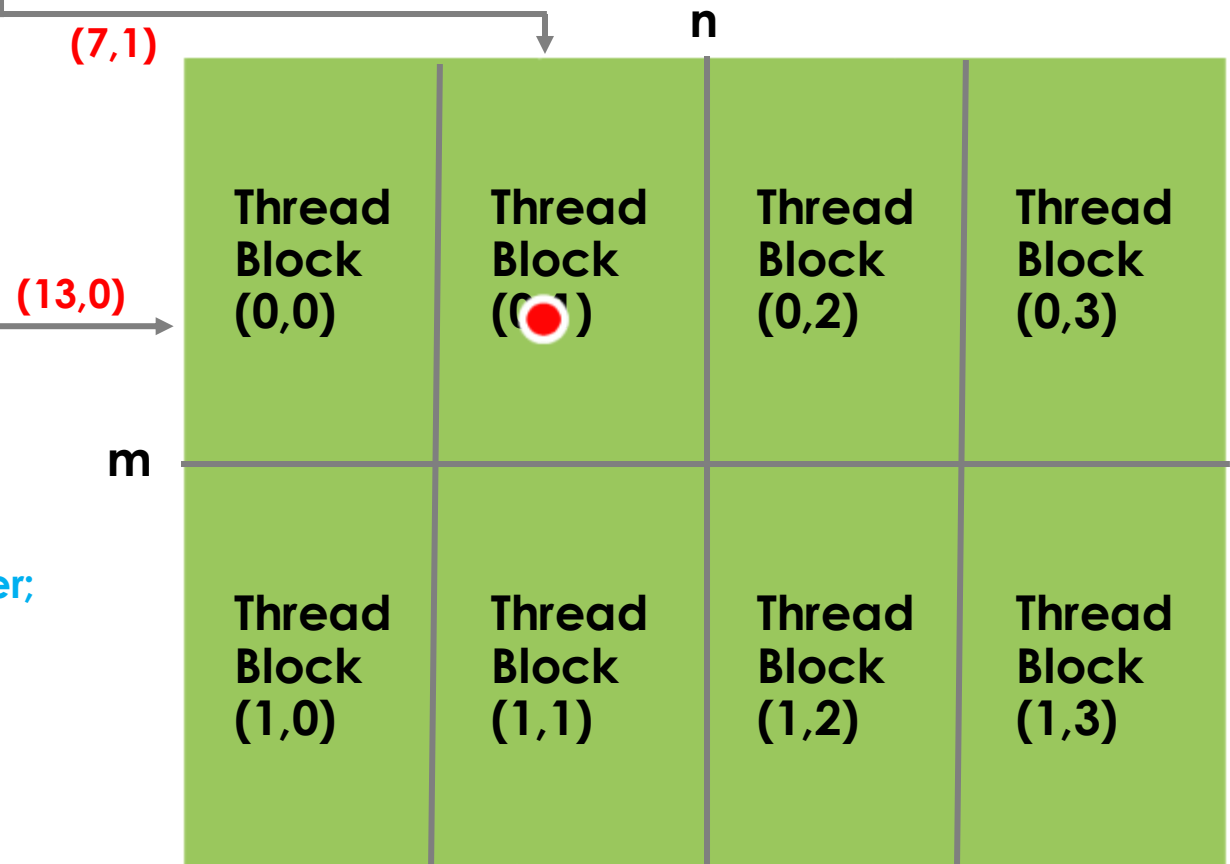
```

Linear offsets are used for addressing A and B storage

Load element of A;  
Load element of B;  
Multiply;  
Accumulate into register;

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$



Matrix C(m,n)

# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```

template <typename T>
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
{
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;

    size_t n_pos = ty;
    while(n_pos < n){

        size_t m_pos = tx;
        while(m_pos < m){

            T tmp = static_cast<T>(0.0);
            for(size_t k_pos = 0; k_pos < k; ++k_pos){
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];
            }
            dest[n_pos*m + m_pos] += tmp;

            m_pos += blockDim.x;
        }

        n_pos += blockDim.y;
    }
    return;
}

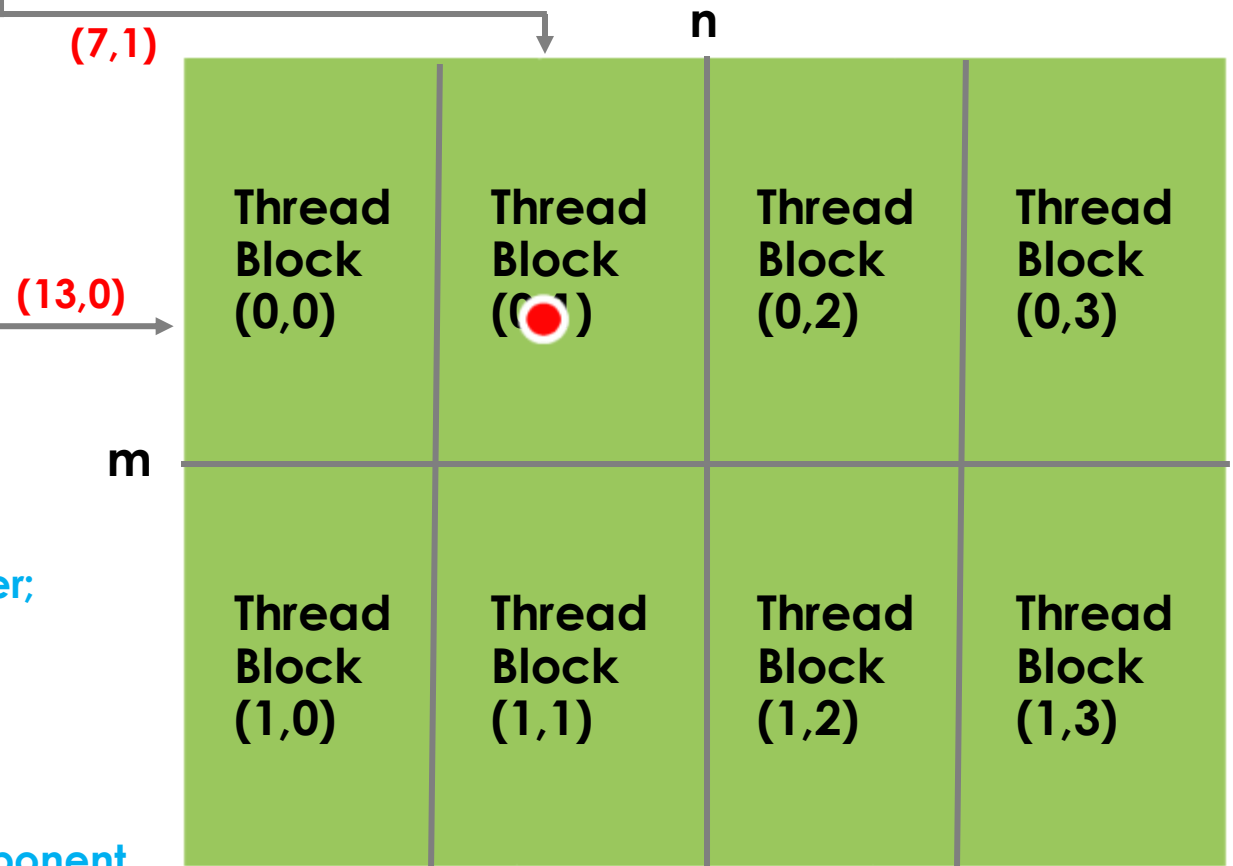
```

Each CUDA thread block computes:  
 $C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

Linear offsets are used for addressing A and B storage

Load element of A;  
 Load element of B;  
 Multiply;  
 Accumulate into register;

Global memory accesses to A and B are coalesced: threadIdx.x is the minor component



Matrix C(m,n)

# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```
template <typename T>
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
{
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;

    size_t n_pos = ty;
    while(n_pos < n){

        size_t m_pos = tx;
        while(m_pos < m){

            T tmp = static_cast<T>(0.0);
            for(size_t k_pos = 0; k_pos < k; ++k_pos){
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];
            }
            dest[n_pos*m + m_pos] += tmp;
            Upload register to global memory

            m_pos += blockDim.x*blockDim.x;
        }

        n_pos += blockDim.y*blockDim.y;
    }
    return;
}
```

Each CUDA thread block computes:  
 $C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

The diagram illustrates the layout of Matrix C(m,n) as a 2x4 grid of thread blocks. The vertical axis is labeled 'm' and the horizontal axis is labeled 'n'. The thread blocks are labeled as follows:

- Row 0: Thread Block (0,0), Thread Block (0,1), Thread Block (0,2), Thread Block (0,3)
- Row 1: Thread Block (1,0), Thread Block (1,1), Thread Block (1,2), Thread Block (1,3)

Arrows indicate the mapping of thread indices to block coordinates:

- An arrow from `ty` points to the vertical axis, with the label  $(7,1)$  indicating the row index.
- An arrow from `tx` points to the horizontal axis, with the label  $(13,0)$  indicating the column index.

The thread block at (0,1) is highlighted with a red dot, indicating the current thread's position.

# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```
template <typename T>
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
{
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;

    size_t n_pos = ty;
    while(n_pos < n){

        size_t m_pos = tx;
        while(m_pos < m){

            T tmp = static_cast<T>(0.0);
            for(size_t k_pos = 0; k_pos < k; ++k_pos){
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];
            }
            dest[n_pos*m + m_pos] += tmp;

            m_pos += blockDim.x*blockDim.x;
        }

        n_pos += blockDim.y*blockDim.y;
    }
    return;
}
```

Each CUDA thread block computes:  
 $C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

(7,1)

(13,0)

m

n

Thread Block (0,0)

Thread Block (0,1)

Thread Block (0,2)

Thread Block (0,3)

Thread Block (1,0)

Thread Block (1,1)

Thread Block (1,2)

Thread Block (1,3)

Matrix C(m,n)

Loop over X and Y dims of C in case CUDA thread blocks do not cover full matrix C

# CUDA BLA Library: Naïve GEMM kernel (algorithm 0)

```
template <typename T>
__global__ void gpu_gemm_nn(int m, int n, int k, T * __restrict__ dest, const T * __restrict__ left, const T * __restrict__ right)
{
    size_t ty = blockIdx.y*blockDim.y + threadIdx.y;
    size_t tx = blockIdx.x*blockDim.x + threadIdx.x;

    size_t n_pos = ty;
    while(n_pos < n){

        size_t m_pos = tx;
        while(m_pos < m){

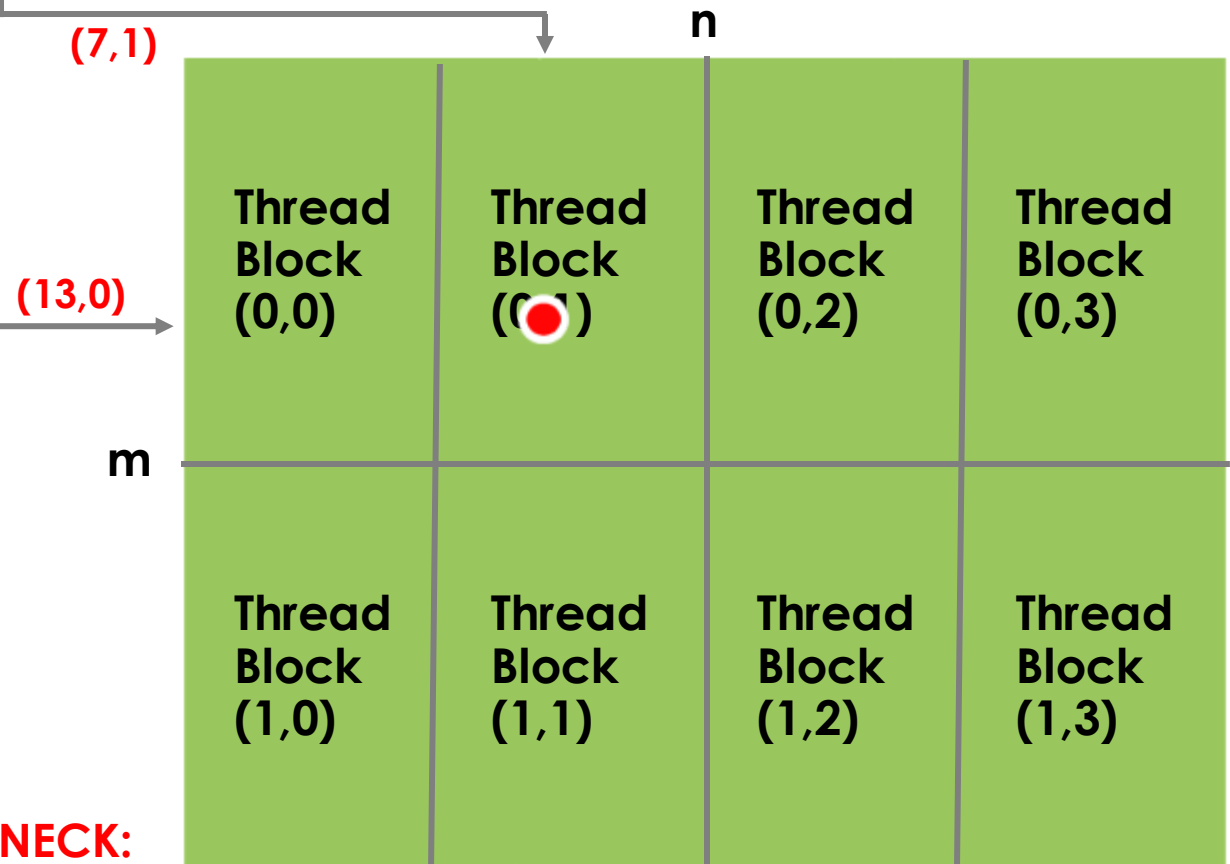
            T tmp = static_cast<T>(0.0);
            for(size_t k_pos = 0; k_pos < k; ++k_pos){
                tmp += left[k_pos*m + m_pos] * right[n_pos*k + k_pos];
            }
            dest[n_pos*m + m_pos] += tmp;

            m_pos += blockDim.x*blockDim.x;
        }

        n_pos += blockDim.y*blockDim.y;
    }
    return;
}
```

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$



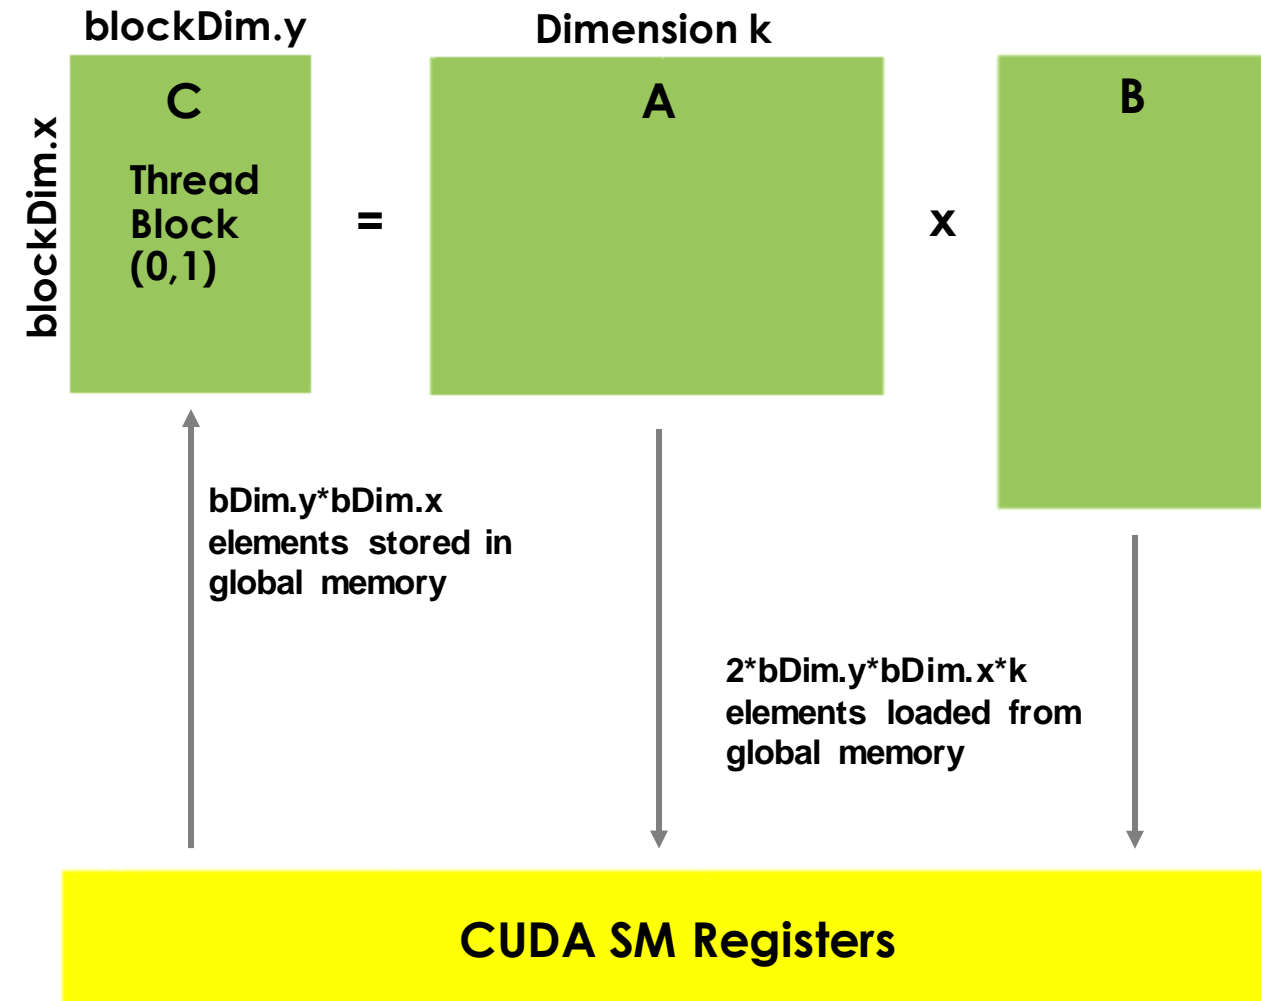
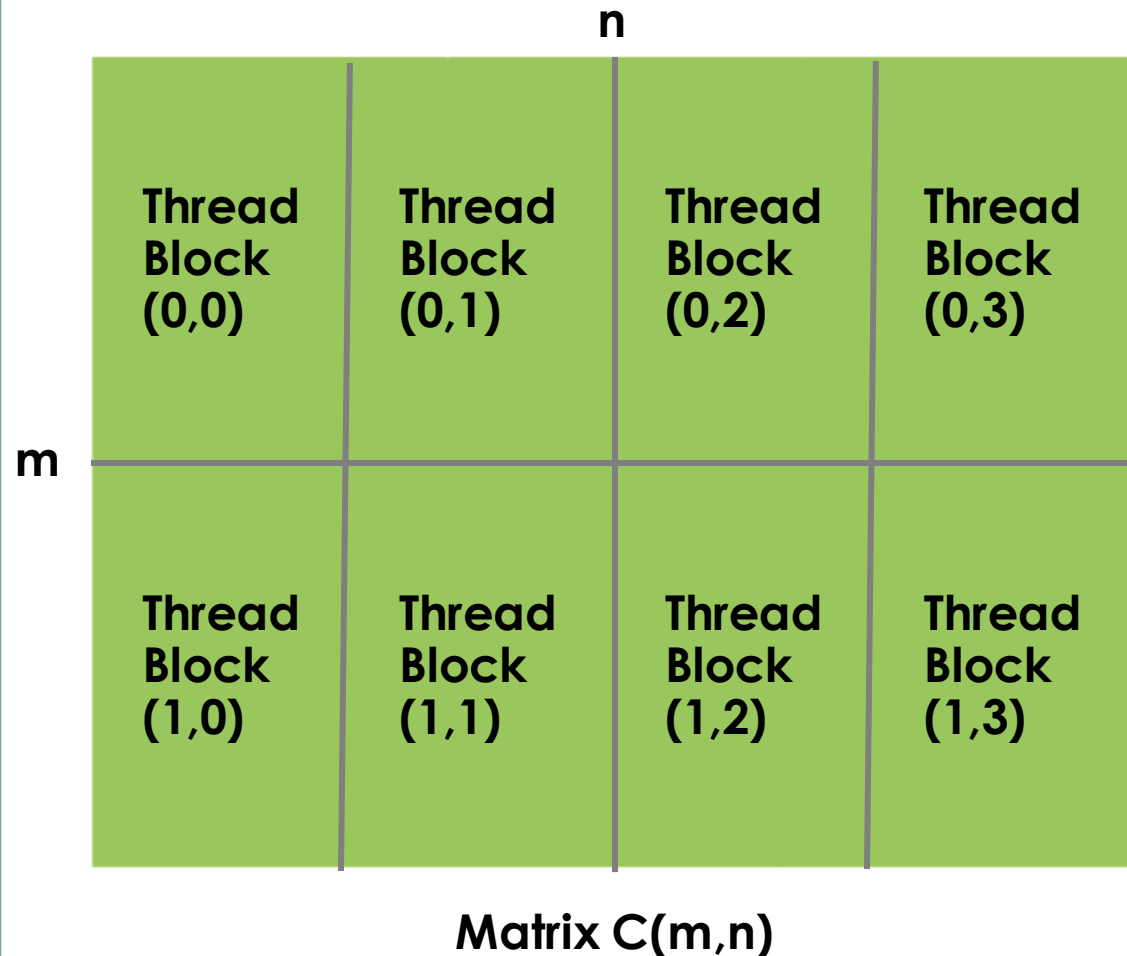
**GLOBAL MEMORY ACCESS BOTTLENECK:**  
 $2 * \text{bDim.y} * \text{bDim.x} * k$  loads per block  
 $2 * m * n * k$  loads per kernel

# CUDA BLA Library: Shared Memory GEMM need

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

$\sim 2 * m * n * k / \text{bDim.x}$  global loads per kernel

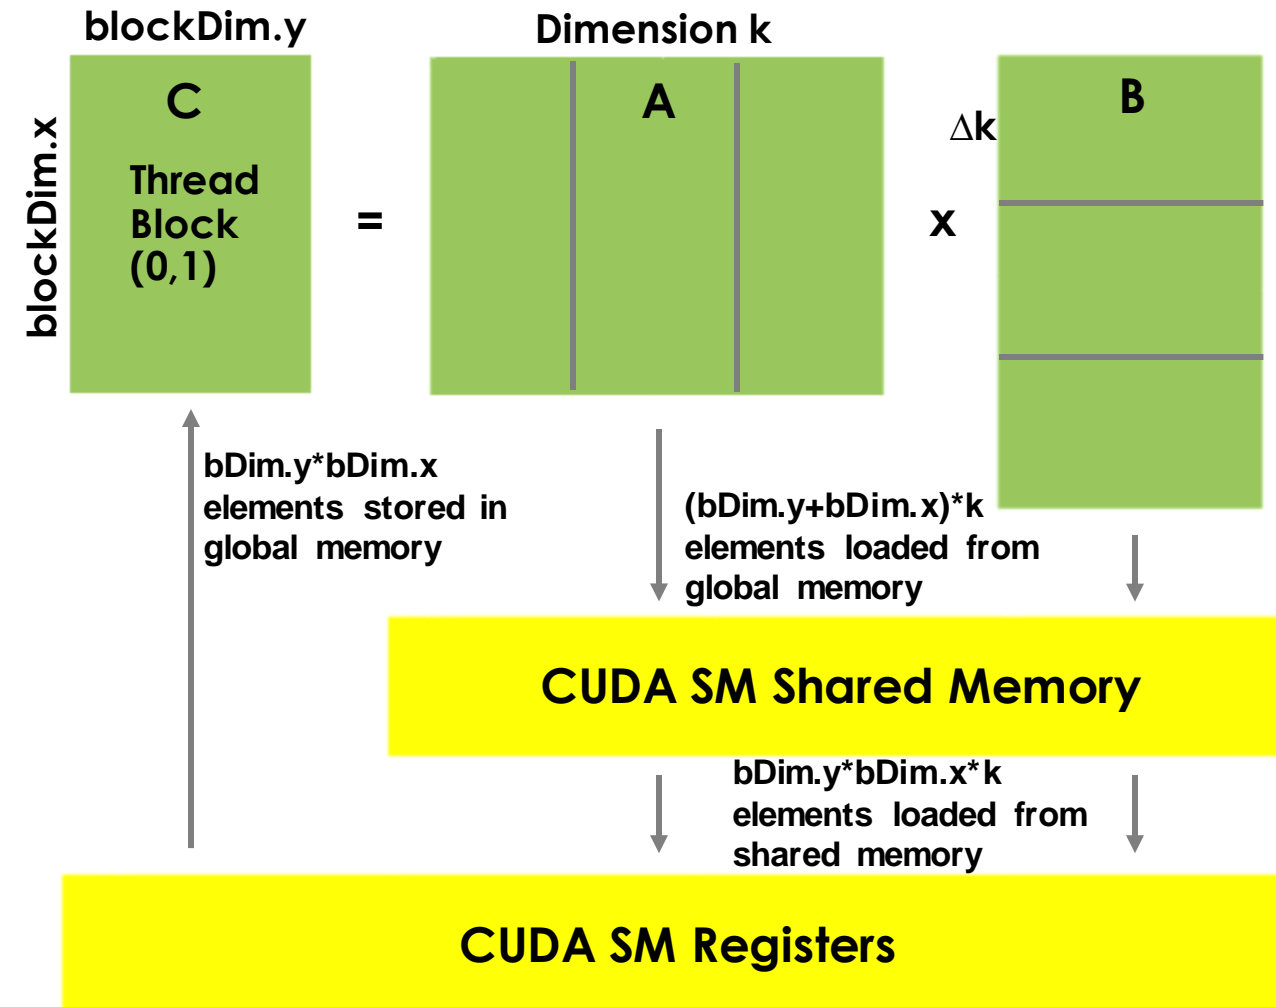
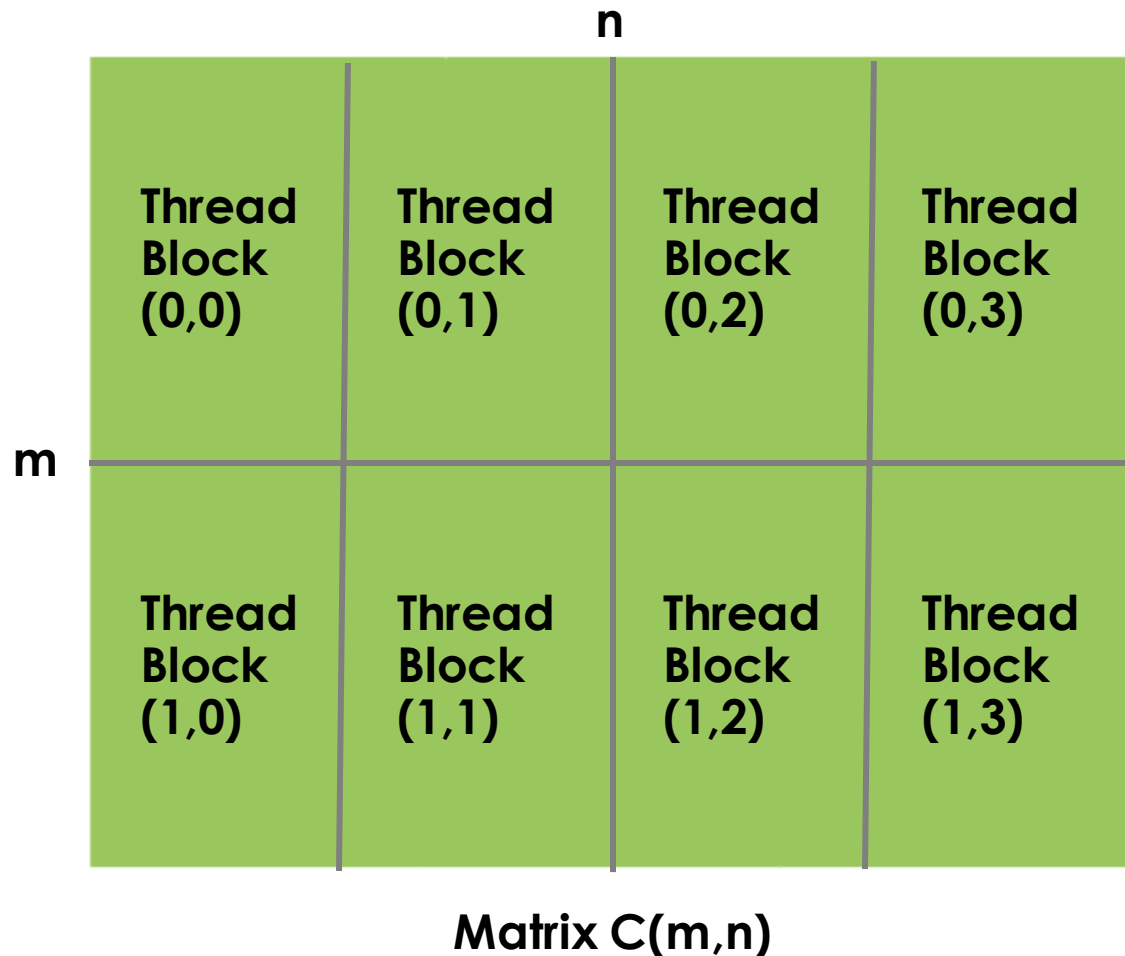


# CUDA BLA Library: Shared Memory GEMM (algorithm 1)

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

$\sim 2 * m * n * k / \text{bDim.x}$  global loads per kernel





# CUDA BLA Library: Shared Memory GEMM (algorithm 1)

TileC(M,N), TileA(M,K), TileB(K,N)

template <typename T, int TILE\_EXT\_N, int TILE\_EXT\_M, int TILE\_EXT\_K>

\_\_global\_\_ void **gpu\_gemm\_sh\_nn**(int m, int n, int k, //in: matrix dimensions: C(m,n)+=A(m,k)\*B(k,n)

T \* \_\_restrict\_\_ dest, //inout: pointer to C matrix data

const T \* \_\_restrict\_\_ left, //in: pointer to A matrix data

const T \* \_\_restrict\_\_ right) //in: pointer to B matrix data

{  
using int\_t = int; //either int or size\_t **Shared memory buffers (per thread block)**  
\_\_shared\_\_ T lbuf[TILE\_EXT\_K][TILE\_EXT\_M], rbuf[TILE\_EXT\_N][TILE\_EXT\_K];

for(int\_t n\_pos = blockIdx.y\*blockDim.y; n\_pos < n; n\_pos += gridDim.y\*blockDim.y){ //tile offset in Y

for(int\_t m\_pos = blockIdx.x\*blockDim.x; m\_pos < m; m\_pos += gridDim.x\*blockDim.x){ //tile offset in X

T tmp = static\_cast<T>(0.0); //accumulator

for(int\_t k\_pos = 0; k\_pos < k; k\_pos += TILE\_EXT\_K){ //tile begin position along dimension K  
int\_t k\_end = k\_pos + TILE\_EXT\_K; if(k\_end > k) k\_end = k;

//Load a tile of matrix A(m\_pos:TILE\_EXT\_M, k\_pos:TILE\_EXT\_K):

if(m\_pos + threadIdx.x < m){  
for(int\_t k\_loc = k\_pos + threadIdx.y; k\_loc < k\_end; k\_loc += blockDim.y){  
lbuf[k\_loc-k\_pos][threadIdx.x] = left[k\_loc\*m + (m\_pos+threadIdx.x)];  
}  
}

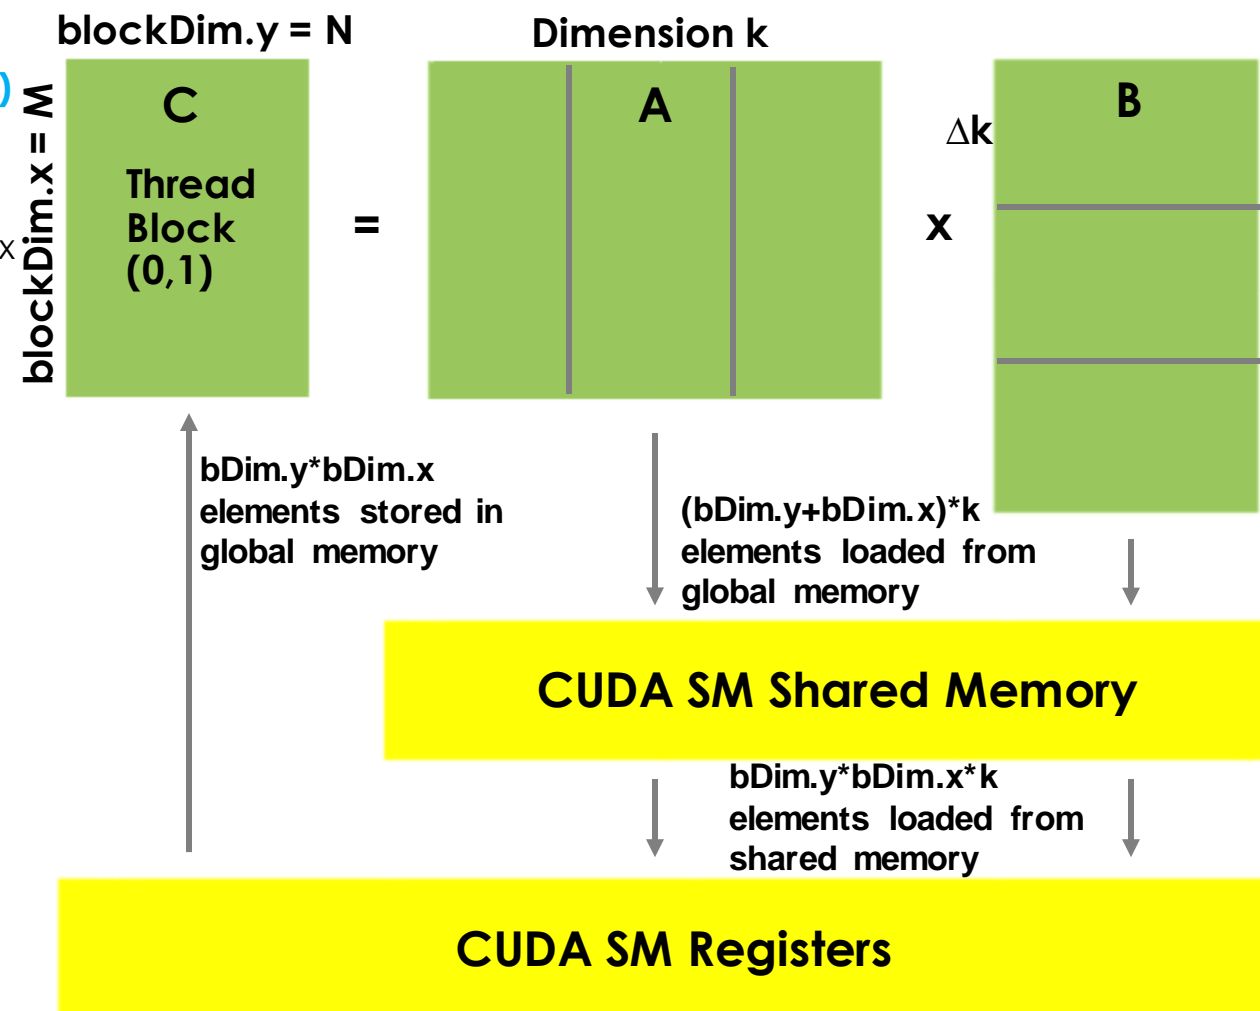
//Load a tile of matrix B(k\_pos:TILE\_EXT\_K, n\_pos:TILE\_EXT\_N):

if(n\_pos + threadIdx.y < n){  
for(int\_t k\_loc = k\_pos + threadIdx.x; k\_loc < k\_end; k\_loc += blockDim.x){  
rbuf[threadIdx.y][k\_loc-k\_pos] = right[(n\_pos+threadIdx.y)\*k + k\_loc];  
}  
}

\_\_syncthreads();

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$



# CUDA BLA Library: Shared Memory GEMM (algorithm 1)

**TileC(M,N), TileA(M,K), TileB(K,N)**

```
template <typename T, int TILE_EXT_N, int TILE_EXT_M, int TILE_EXT_K>
__global__ void gpu_gemm_sh_nn(int m, int n, int k, //in: matrix dimensions: C(m,n)+=A(m,k)*B(k,n)
    T * __restrict__ dest,          //inout: pointer to C matrix data
    const T * __restrict__ left,    //in: pointer to A matrix data
    const T * __restrict__ right) //in: pointer to B matrix data
{
    using int_t = int; //either int or size_t
    __shared__ T lbuf[TILE_EXT_K][TILE_EXT_M], rbuf[TILE_EXT_N][TILE_EXT_K];

    for(int_t n_pos = blockIdx.y*blockDim.y; n_pos < n; n_pos += gridDim.y*blockDim.y){ //tile offset in Y

        for(int_t m_pos = blockIdx.x*blockDim.x; m_pos < m; m_pos += gridDim.x*blockDim.x){ //tile offset in X

            T tmp = static_cast<T>(0.0); //accumulator

            for(int_t k_pos = 0; k_pos < k; k_pos += TILE_EXT_K){ //tile begin position along dimension K
                int_t k_end = k_pos + TILE_EXT_K; if(k_end > k) k_end = k;

                //Load a tile of matrix A(m_pos:TILE_EXT_M, k_pos:TILE_EXT_K):
                if(m_pos + threadIdx.x < m){
                    for(int_t k_loc = k_pos + threadIdx.y; k_loc < k_end; k_loc += blockDim.y){
                        lbuf[k_loc-k_pos][threadIdx.x] = left[k_loc*m + (m_pos+threadIdx.x)];
                    }
                }

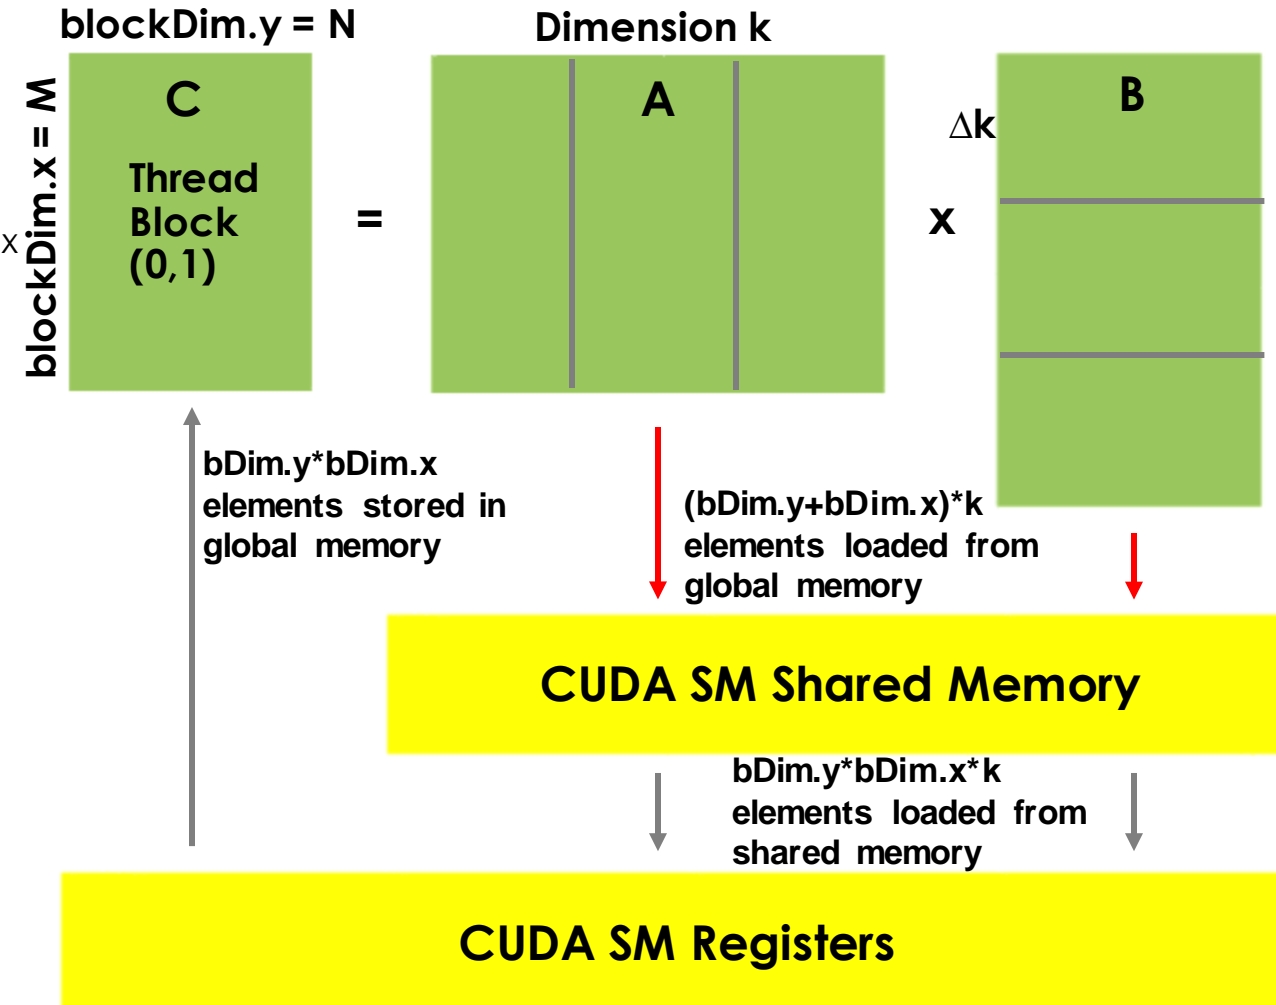
                //Load a tile of matrix B(k_pos:TILE_EXT_K, n_pos:TILE_EXT_N):
                if(n_pos + threadIdx.y < n){
                    for(int_t k_loc = k_pos + threadIdx.x; k_loc < k_end; k_loc += blockDim.x){
                        rbuf[threadIdx.y][k_loc-k_pos] = right[(n_pos+threadIdx.y)*k + k_loc];
                    }
                }

                __syncthreads();
            }
        }
    }
}
```

**Loading shared memory buffers**

**Each CUDA thread block computes:**

**$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$**



# CUDA BLA Library: Shared Memory GEMM (algorithm 1)

**TileC(M,N), TileA(M,K), TileB(K,N)**

```
template <typename T, int TILE_EXT_N, int TILE_EXT_M, int TILE_EXT_K>
__global__ void gpu_gemm_sh_nn(int m, int n, int k, //in: matrix dimensions: C(m,n)+=A(m,k)*B(k,n)
    T * __restrict__ dest, //inout: pointer to C matrix data
    const T * __restrict__ left, //in: pointer to A matrix data
    const T * __restrict__ right) //in: pointer to B matrix data
{
    using int_t = int; //either int or size_t
    __shared__ T lbuf[TILE_EXT_K][TILE_EXT_M], rbuf[TILE_EXT_N][TILE_EXT_K];

    for(int_t n_pos = blockIdx.y*blockDim.y; n_pos < n; n_pos += gridDim.y*blockDim.y){ //tile offset in Y

        for(int_t m_pos = blockIdx.x*blockDim.x; m_pos < m; m_pos += gridDim.x*blockDim.x){ //tile offset in X

            T tmp = static_cast<T>(0.0); //accumulator

            for(int_t k_pos = 0; k_pos < k; k_pos += TILE_EXT_K){ //tile begin position along dimension K
                int_t k_end = k_pos + TILE_EXT_K; if(k_end > k) k_end = k;

                //Load a tile of matrix A(m_pos:TILE_EXT_M, k_pos:TILE_EXT_K):
                if(m_pos + threadIdx.x < m){
                    for(int_t k_loc = k_pos + threadIdx.y; k_loc < k_end; k_loc += blockDim.y){
                        lbuf[k_loc-k_pos][threadIdx.x] = left[k_loc*m + (m_pos+threadIdx.x)];
                    }
                }

                //Load a tile of matrix B(k_pos:TILE_EXT_K, n_pos:TILE_EXT_N):
                if(n_pos + threadIdx.y < n){
                    for(int_t k_loc = k_pos + threadIdx.x; k_loc < k_end; k_loc += blockDim.x){
                        rbuf[threadIdx.y][k_loc-k_pos] = right[(n_pos+threadIdx.y)*k + k_loc];
                    }
                }

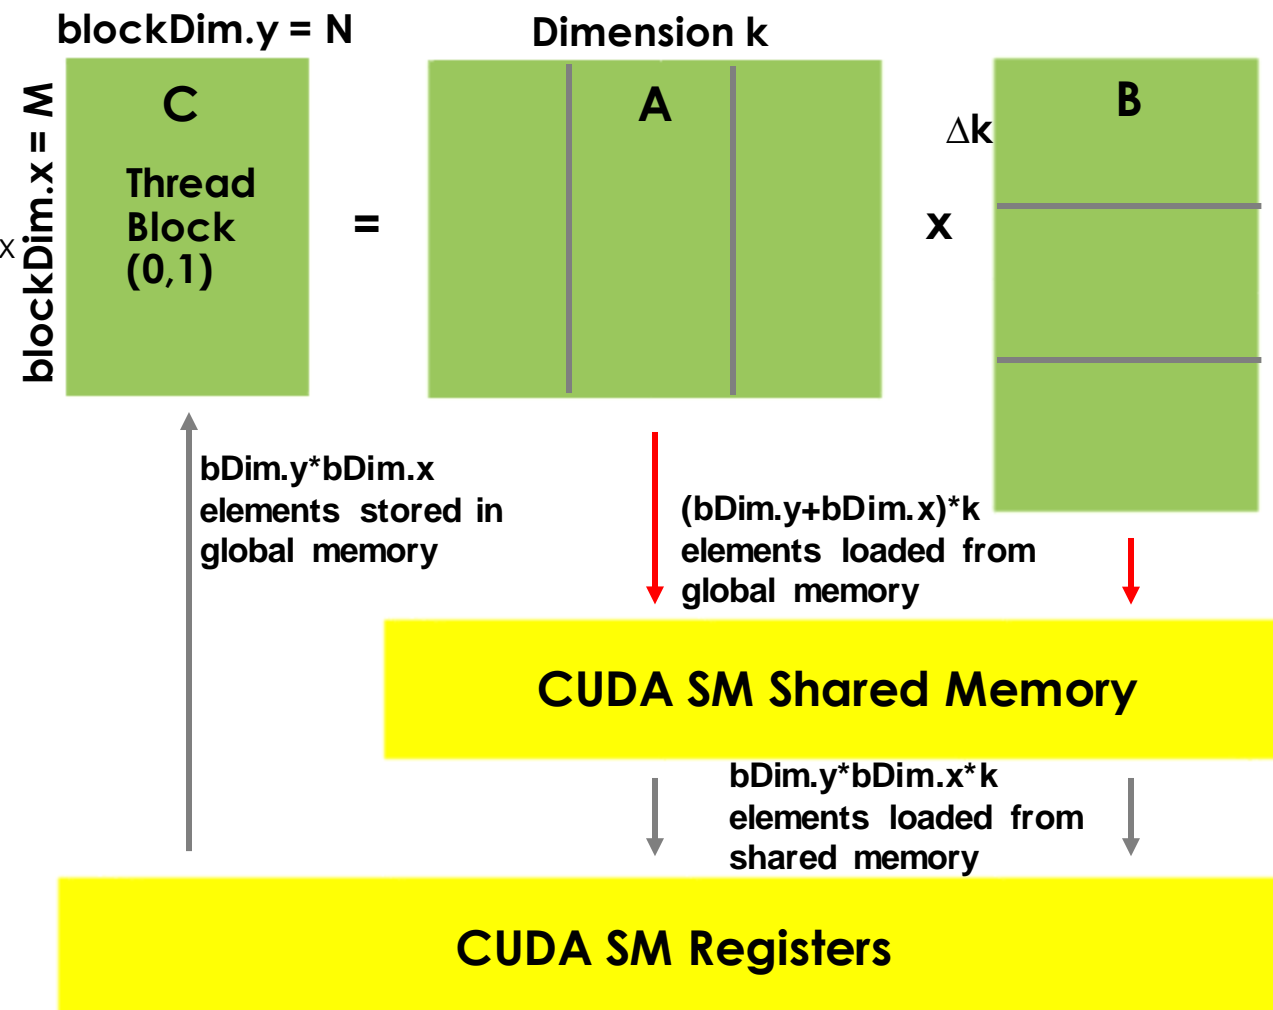
                __syncthreads();
            }
        }
    }
}
```

**Loading shared memory buffers**

**Global memory accesses to A and B are coalesced: threadIdx.x is the minor component**

**Each CUDA thread block computes:**

**$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$**



# CUDA BLA Library: Shared Memory GEMM (algorithm 1)

**TileC(M,N), TileA(M,K), TileB(K,N)**

```
template <typename T, int TILE_EXT_N, int TILE_EXT_M, int TILE_EXT_K>
__global__ void gpu_gemm_sh_nn(int m, int n, int k, //in: matrix dimensions: C(m,n)+=A(m,k)*B(k,n)
    T * __restrict__ dest,          //inout: pointer to C matrix data
    const T * __restrict__ left,    //in: pointer to A matrix data
    const T * __restrict__ right) //in: pointer to B matrix data
{
    using int_t = int; //either int or size_t
    __shared__ T lbuf[TILE_EXT_K][TILE_EXT_M], rbuf[TILE_EXT_N][TILE_EXT_K];

    for(int_t n_pos = blockIdx.y*blockDim.y; n_pos < n; n_pos += gridDim.y*blockDim.y){ //tile offset in Y

        for(int_t m_pos = blockIdx.x*blockDim.x; m_pos < m; m_pos += gridDim.x*blockDim.x){ //tile offset in X

            T tmp = static_cast<T>(0.0); //accumulator

            for(int_t k_pos = 0; k_pos < k; k_pos += TILE_EXT_K){ //tile begin position along dimension K
                int_t k_end = k_pos + TILE_EXT_K; if(k_end > k) k_end = k;

                //Load a tile of matrix A(m_pos:TILE_EXT_M, k_pos:TILE_EXT_K):
                if(m_pos + threadIdx.x < m){
                    for(int_t k_loc = k_pos + threadIdx.y; k_loc < k_end; k_loc += blockDim.y){
                        lbuf[k_loc-k_pos][threadIdx.x] = left[k_loc*m + (m_pos+threadIdx.x)];
                    }
                }

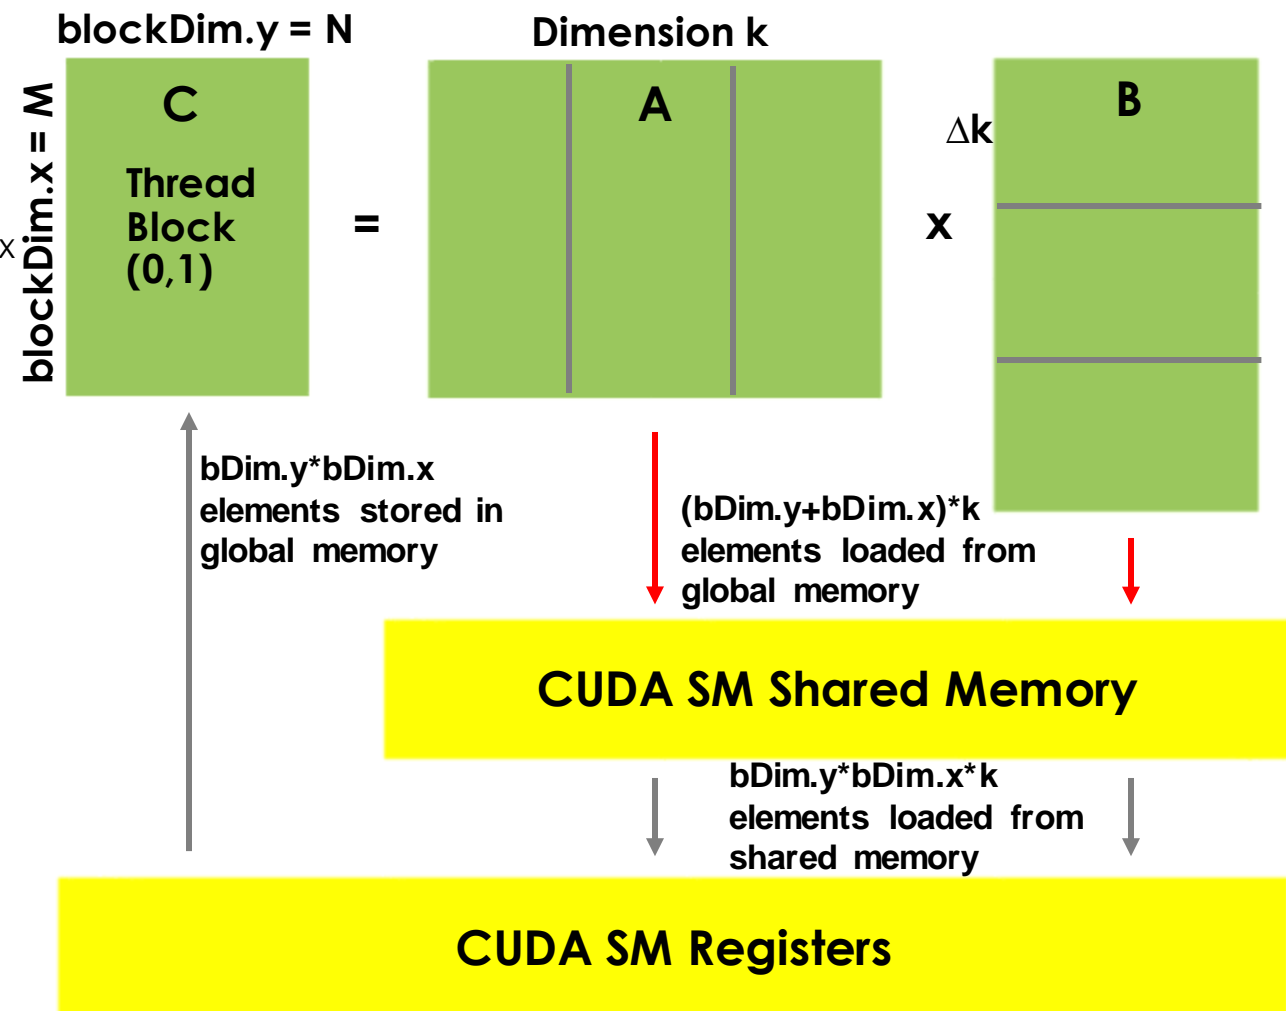
                //Load a tile of matrix B(k_pos:TILE_EXT_K, n_pos:TILE_EXT_N):
                if(n_pos + threadIdx.y < n){
                    for(int_t k_loc = k_pos + threadIdx.x; k_loc < k_end; k_loc += blockDim.x){
                        rbuf[threadIdx.y][k_loc-k_pos] = right[(n_pos+threadIdx.y)*k + k_loc];
                    }
                }

                __syncthreads();
            }
        }
    }
}
```

**Synchronizes threads in a thread block**

**Each CUDA thread block computes:**

**$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$**



# CUDA BLA Library: Shared Memory GEMM (algorithm 1)

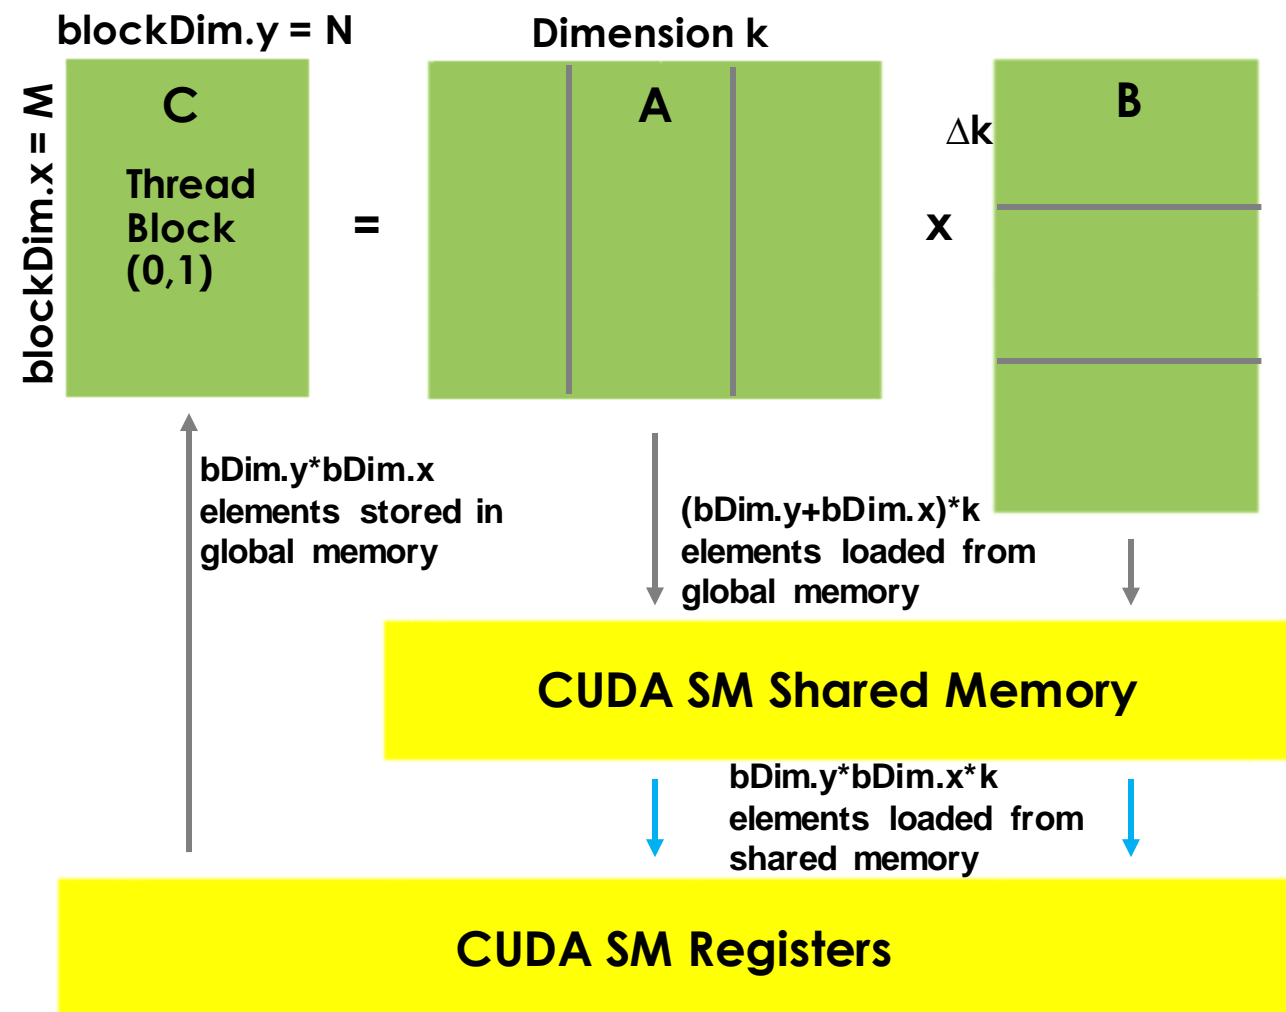
```
//Multiply two loaded tiles to produce a tile of matrix:
if(m_pos + threadIdx.x < m && n_pos + threadIdx.y < n){
    if(k_end - k_pos == TILE_EXT_K){ //known loop count: Unroll
#pragma unroll Unroll loop for performance
        for(int_t l = 0; l < TILE_EXT_K; ++l){
            tmp += lbuf[l][threadIdx.x] * rbuf[threadIdx.y][l];
        }
    }else{ //number of loop iterations is not known at compile time
        for(int_t l = 0; l < (k_end - k_pos); ++l){
            tmp += lbuf[l][threadIdx.x] * rbuf[threadIdx.y][l];
        }
        Performing matrix multiplication
        from shared memory buffers
    }
}
__syncthreads();

} //k_pos

//Store element of the C matrix in global memory:
if(m_pos + threadIdx.x < m && n_pos + threadIdx.y < n)
    dest[(n_pos+threadIdx.y)*m + (m_pos+threadIdx.x)] += tmp;
```

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$



# CUDA BLA Library: Shared Memory GEMM (algorithm 1)

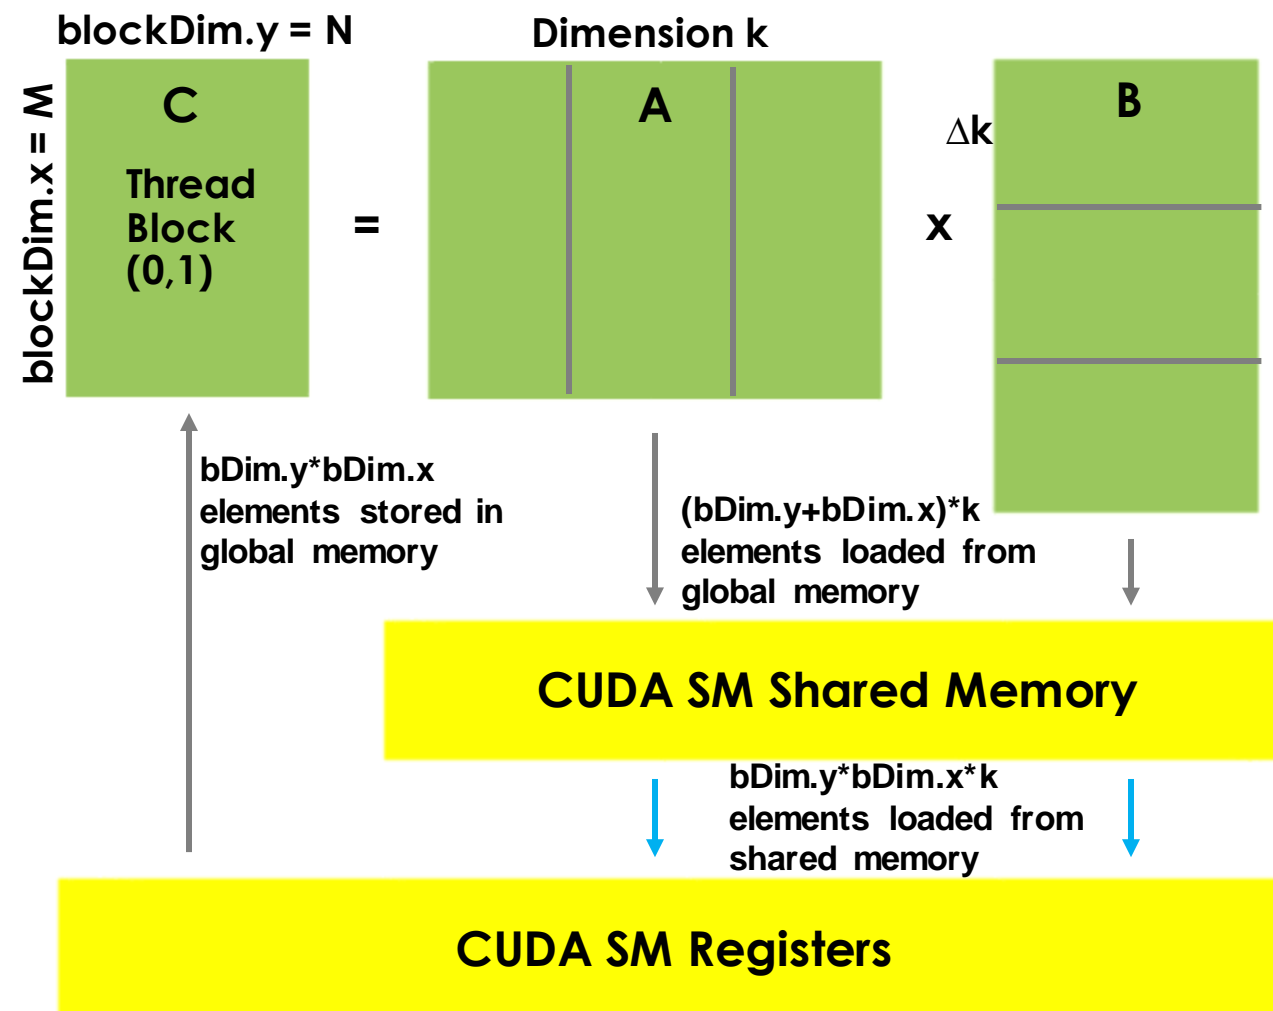
```
//Multiply two loaded tiles to produce a tile of matrix:
if(m_pos + threadIdx.x < m && n_pos + threadIdx.y < n){
    if(k_end - k_pos == TILE_EXT_K){ //known loop count: Unroll
#pragma unroll Unroll loop for performance
        for(int_t l = 0; l < TILE_EXT_K; ++l){
            tmp += lbuf[l][threadIdx.x] * rbuf[threadIdx.y][l];
        }
    }else{ //number of loop iterations is not known at compile time
        for(int_t l = 0; l < (k_end - k_pos); ++l){
            tmp += lbuf[l][threadIdx.x] * rbuf[threadIdx.y][l];
        }
        Performing matrix multiplication
        from shared memory buffers
    }
}
__syncthreads(); Synchronizes threads in a thread block

} //k_pos

//Store element of the C matrix in global memory:
if(m_pos + threadIdx.x < m && n_pos + threadIdx.y < n)
    dest[(n_pos+threadIdx.y)*m + (m_pos+threadIdx.x)] += tmp;
Upload register to global memory (as before)
```

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

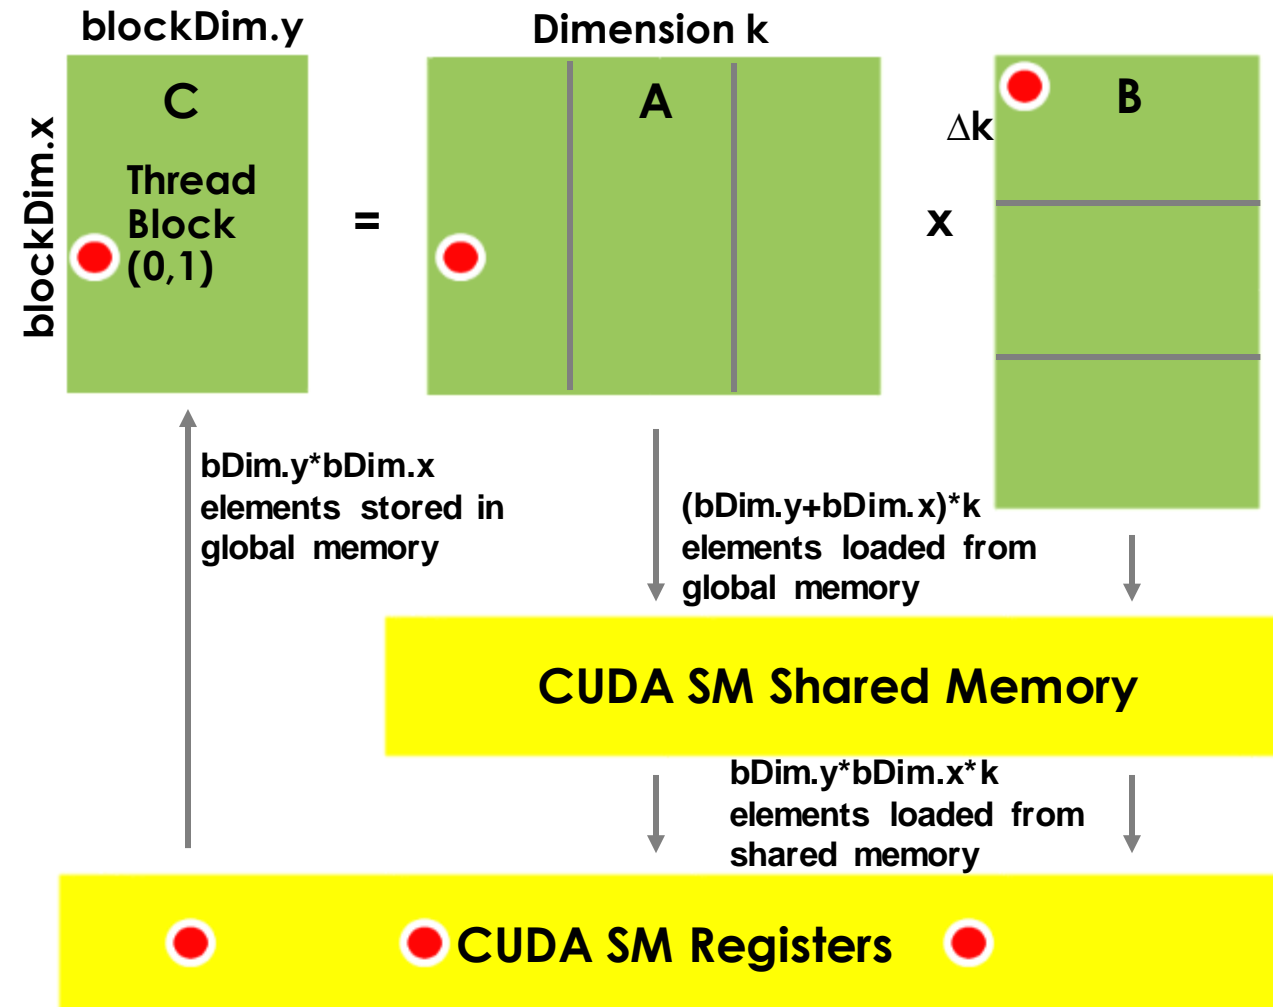
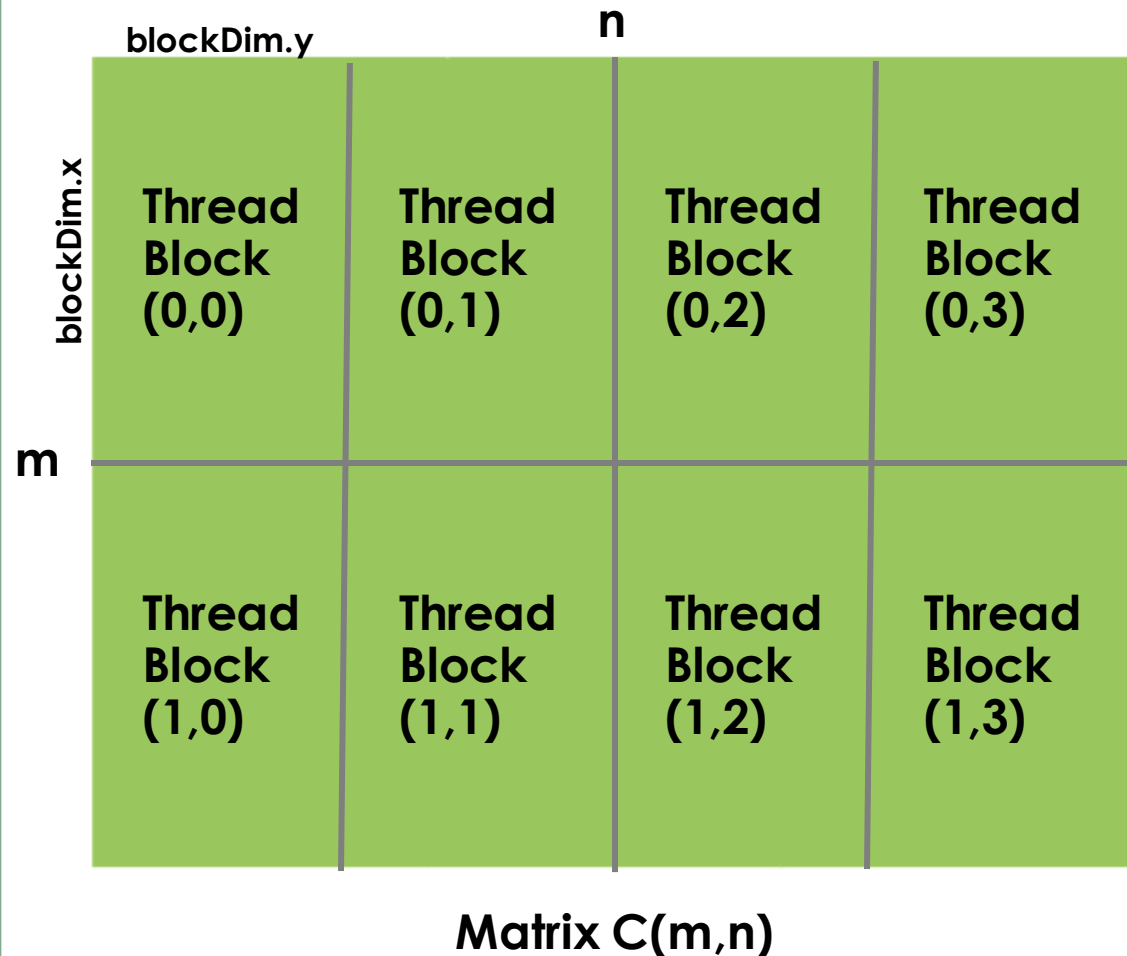


# CUDA BLA Library: +Registers GEMM need

Each CUDA thread block computes:

$C(\text{blockDim.x}, \text{blockDim.y}) += A(\text{blockDim.x}, k) * B(k, \text{blockDim.y})$

$\sim 2 * m * n * k / \text{bDim.x}$  global loads per kernel



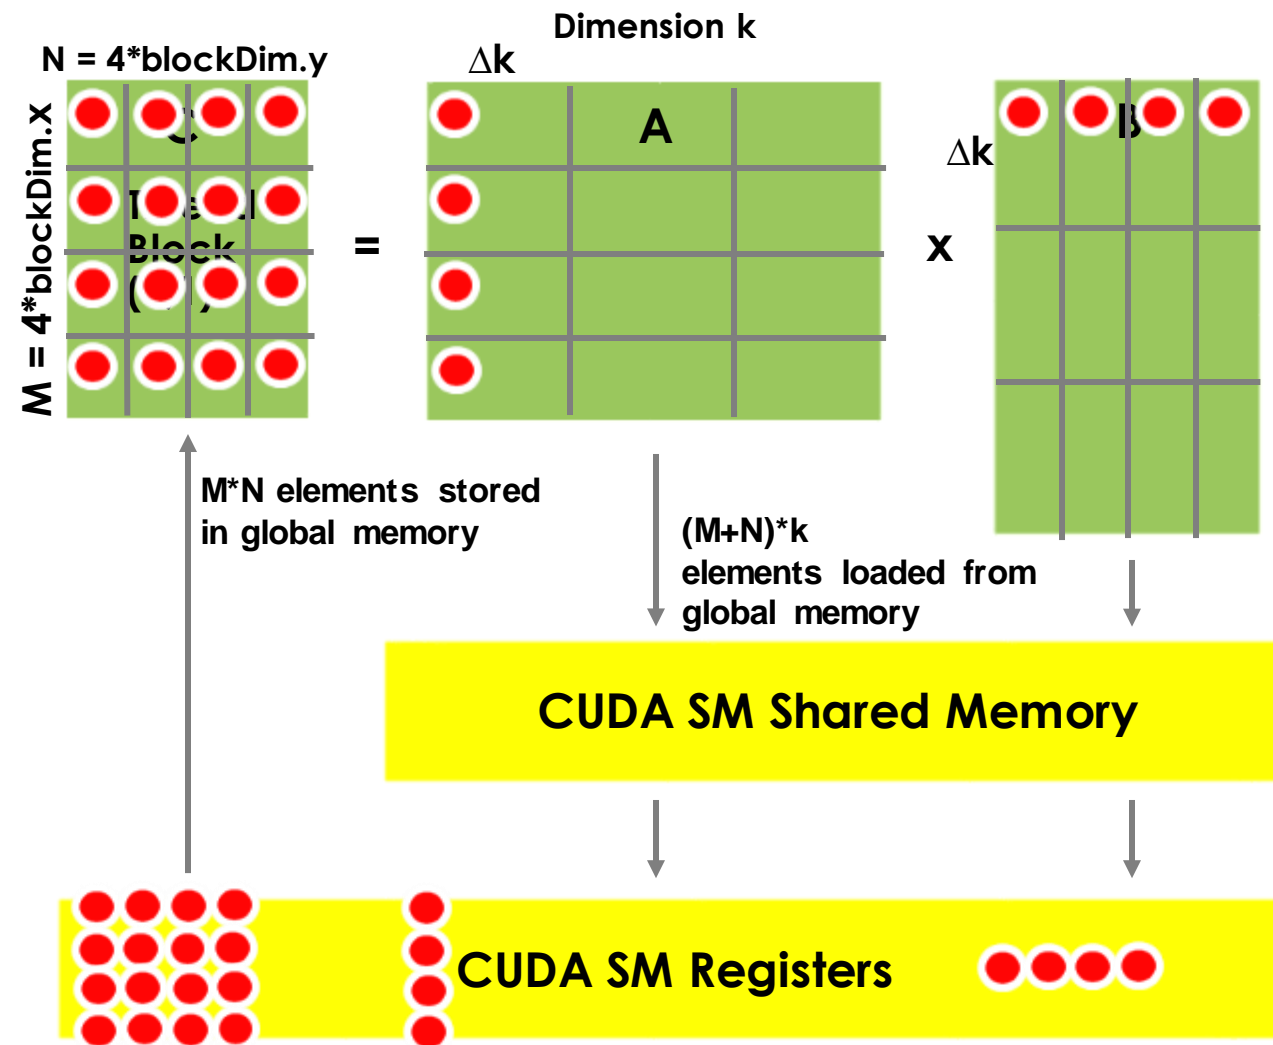
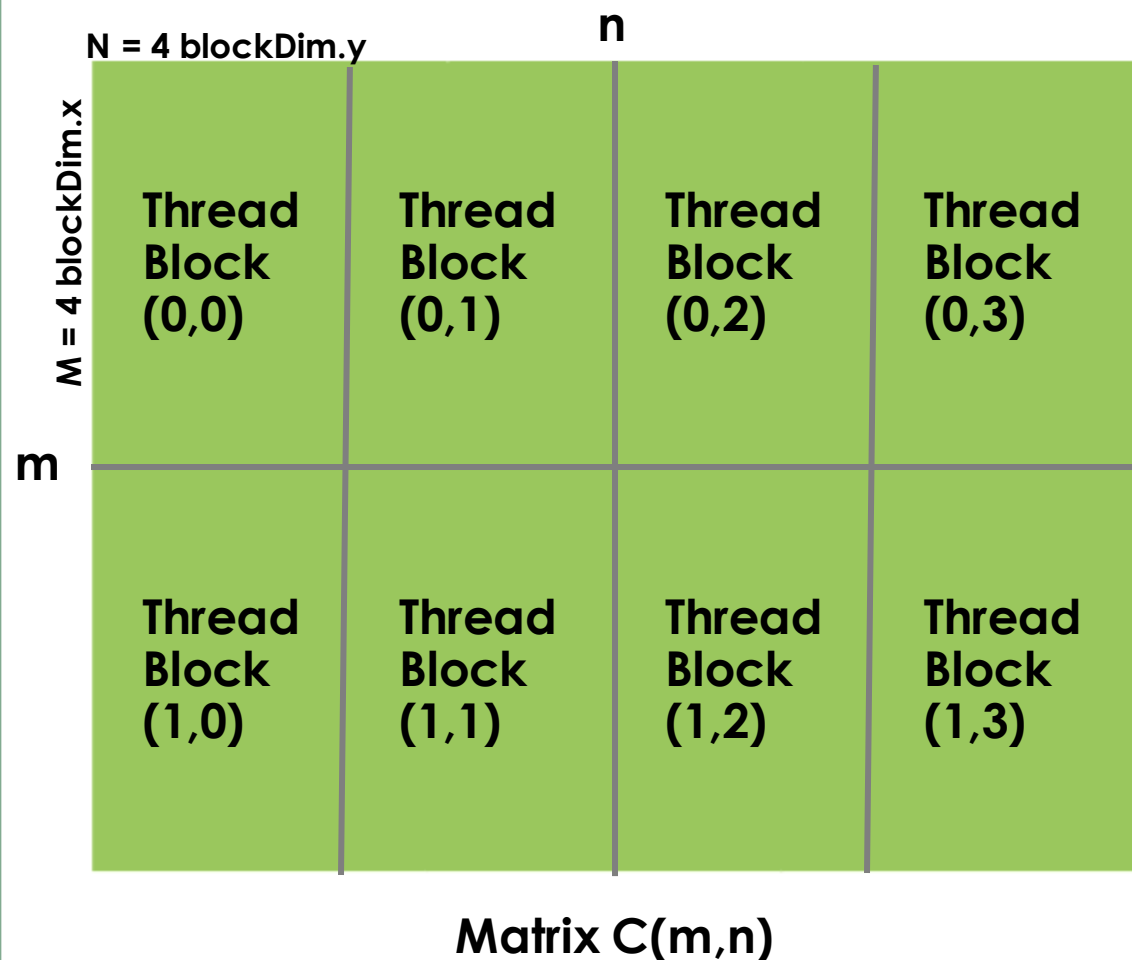


# CUDA BLA Library: +Registers GEMM (algorithm 2)

Each CUDA thread block computes:

$C(M = 4 \cdot \text{blockDim.x}, N = 4 \cdot \text{blockDim.y}) += A(M, k) * B(k, N)$

$\sim 2 \cdot m \cdot n \cdot k / M$  global loads per kernel

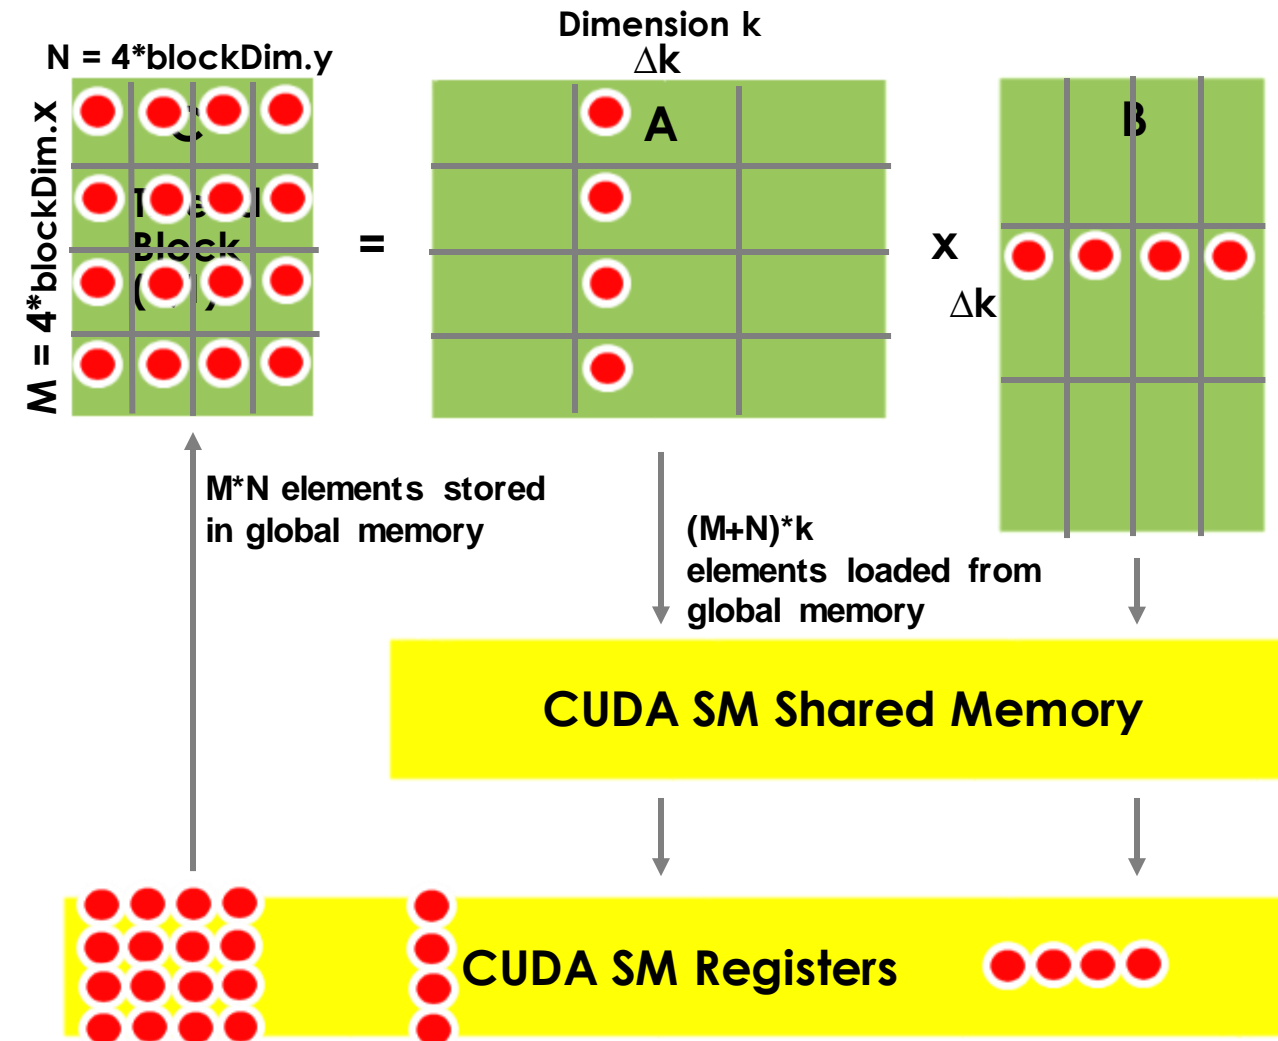
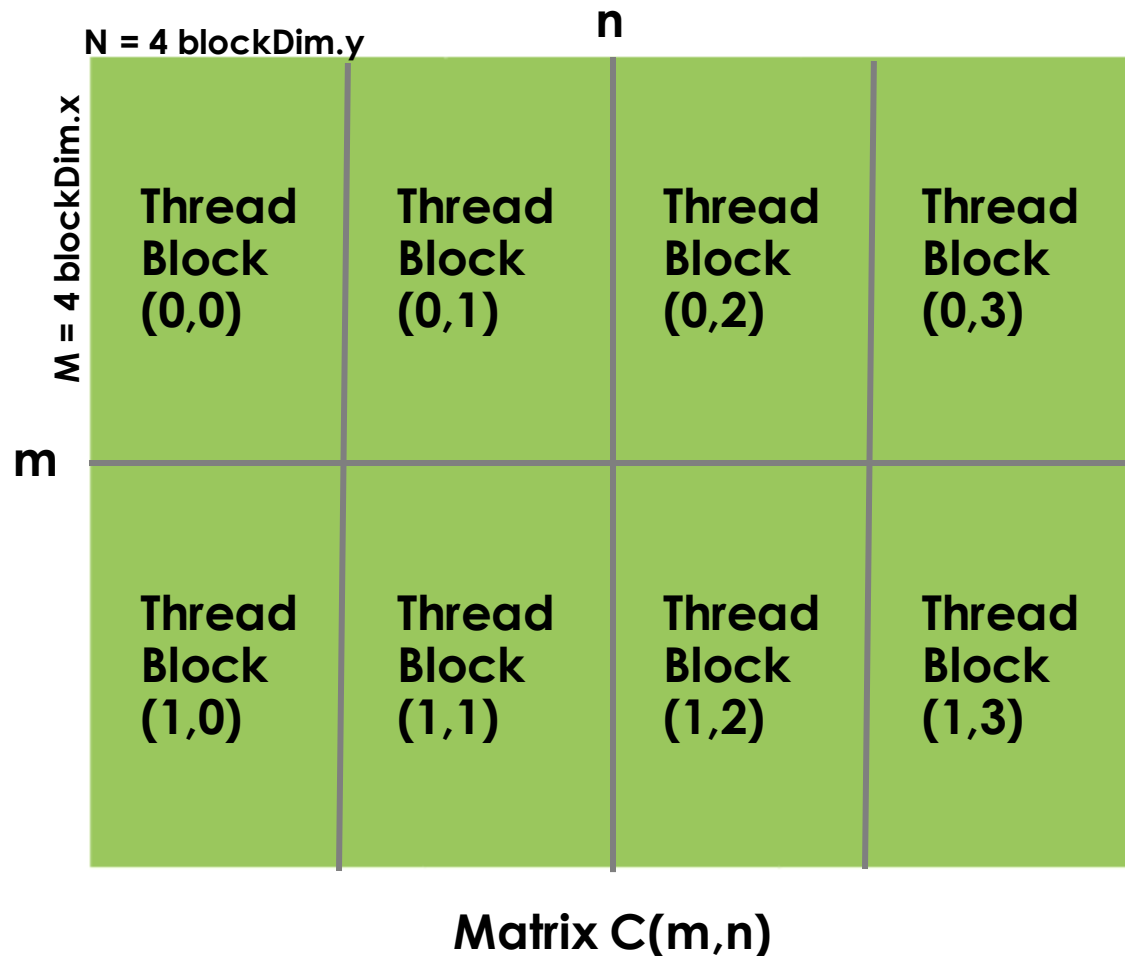


# CUDA BLA Library: +Registers GEMM (algorithm 2)

Each CUDA thread block computes:

$C(M = 4 \cdot \text{blockDim.x}, N = 4 \cdot \text{blockDim.y}) += A(M, k) * B(k, N)$

$\sim 2 \cdot m \cdot n \cdot k / M$  global loads per kernel

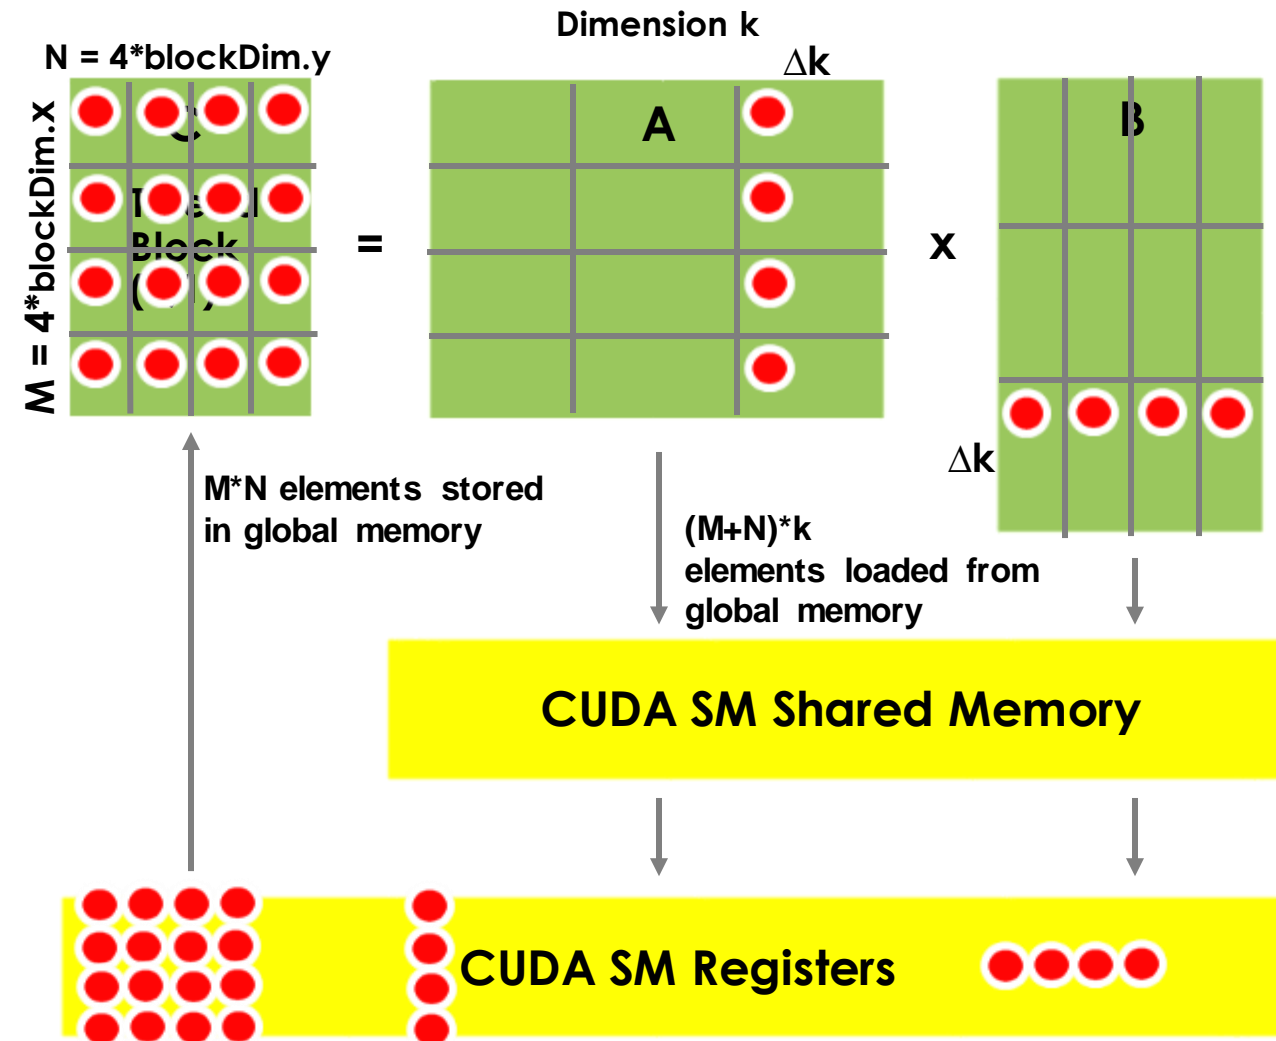
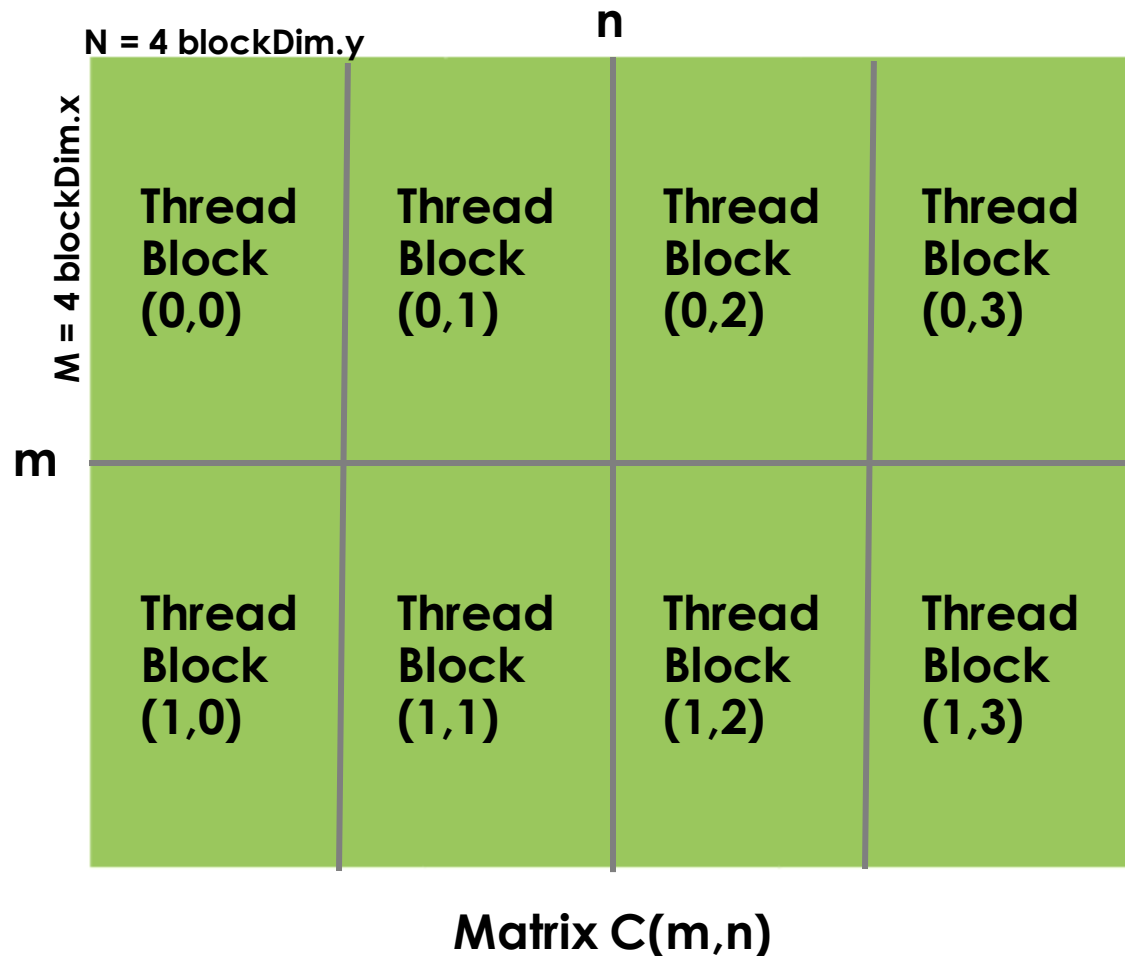


# CUDA BLA Library: +Registers GEMM (algorithm 2)

Each CUDA thread block computes:

$C(M = 4 \cdot \text{blockDim.x}, N = 4 \cdot \text{blockDim.y}) += A(M, k) * B(k, N)$

$\sim 2 \cdot m \cdot n \cdot k / M$  global loads per kernel



# CUDA BLA Library: +Registers GEMM (algorithm 2)

```
template <typename T, int TILE_EXT_N, int TILE_EXT_M, int TILE_EXT_K>
__global__ void gpu_gemm_sh_reg_nn(int m, int n, int k,
    T * __restrict__ dest, //inout: pointer to C matrix data
    const T * __restrict__ left, //in: pointer to A matrix data
    const T * __restrict__ right) //in: pointer to B matrix data
{
    using int_t = int; //either int or size_t
    __shared__ T lbuf[TILE_EXT_K][TILE_EXT_M], rbuf[TILE_EXT_N][TILE_EXT_K];

    for(int_t n_pos = blockIdx.y*TILE_EXT_N; n_pos < n; n_pos += gridDim.y*TILE_EXT_N){
        int_t n_end = n_pos + TILE_EXT_N; if(n_end > n) n_end = n;

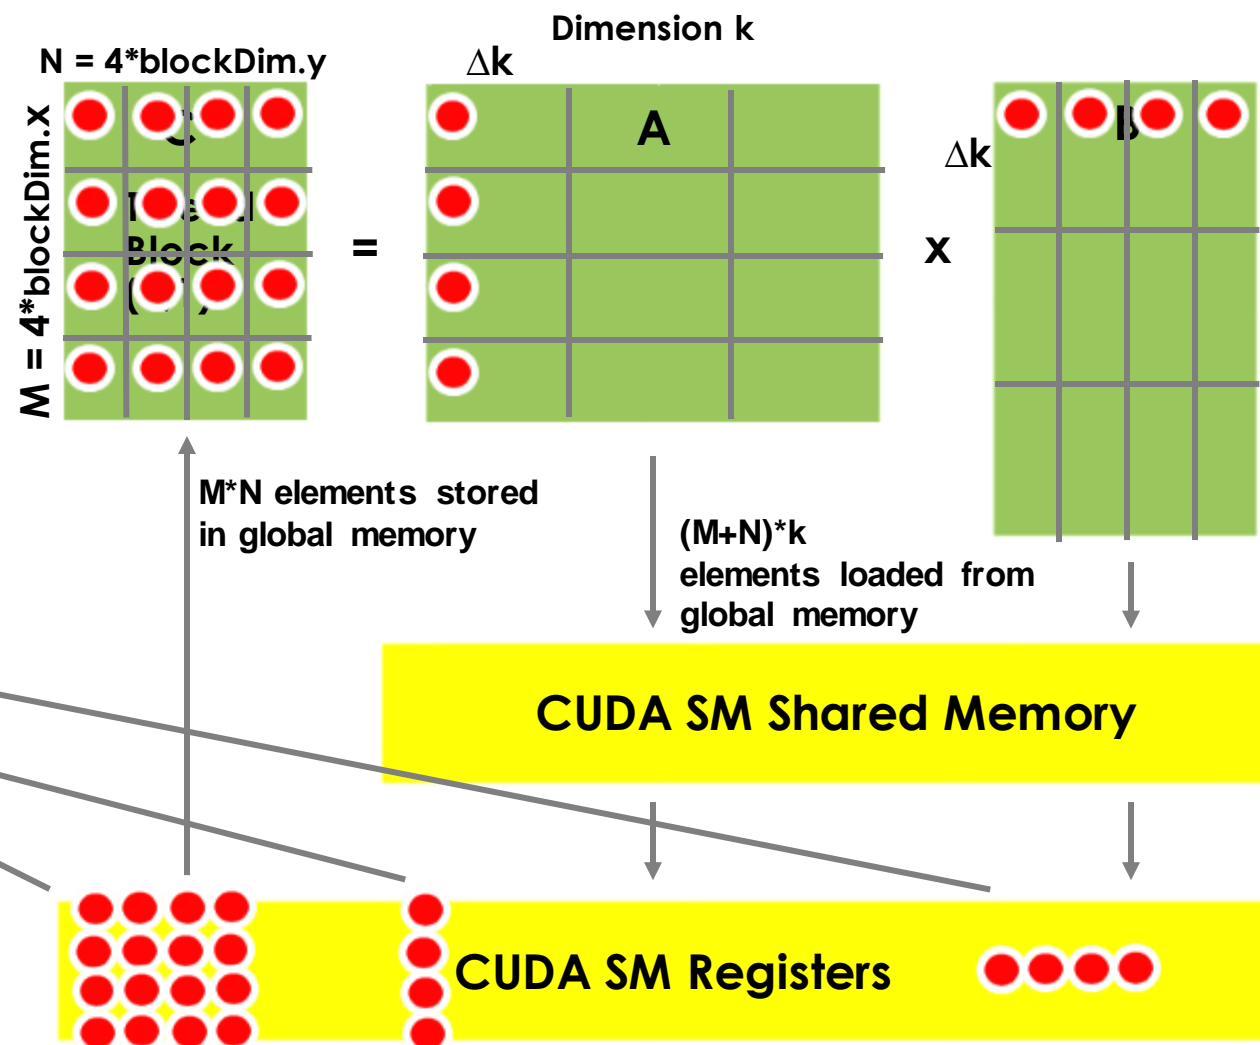
        for(int_t m_pos = blockIdx.x*TILE_EXT_M; m_pos < m; m_pos += gridDim.x*TILE_EXT_M){
            int_t m_end = m_pos + TILE_EXT_M; if(m_end > m) m_end = m;

            if((m_end - m_pos == TILE_EXT_M) && (n_end - n_pos == TILE_EXT_N)){

                //Initialize registers to zero:
                T dreg[4][4] = {static_cast<T>(0.0)};
                T rreg[4] = {static_cast<T>(0.0)};
                T lreg[4] = {static_cast<T>(0.0)};
```

Each CUDA thread block computes:

$C(M = 4*\text{blockDim.x}, N = 4*\text{blockDim.y}) += A(M, k) * B(k, N)$



# CUDA BLA Library: +Registers GEMM (algorithm 2)

```
for(int_t k_pos = 0; k_pos < k; k_pos += TILE_EXT_K){ //k_pos is the position of the CUDA
thread along the K dimension
```

```
    int_t k_end = k_pos + TILE_EXT_K; if(k_end > k) k_end = k;
```

```
    //Load a tile of matrix A(m_pos:TILE_EXT_M, k_pos:TILE_EXT_K):
```

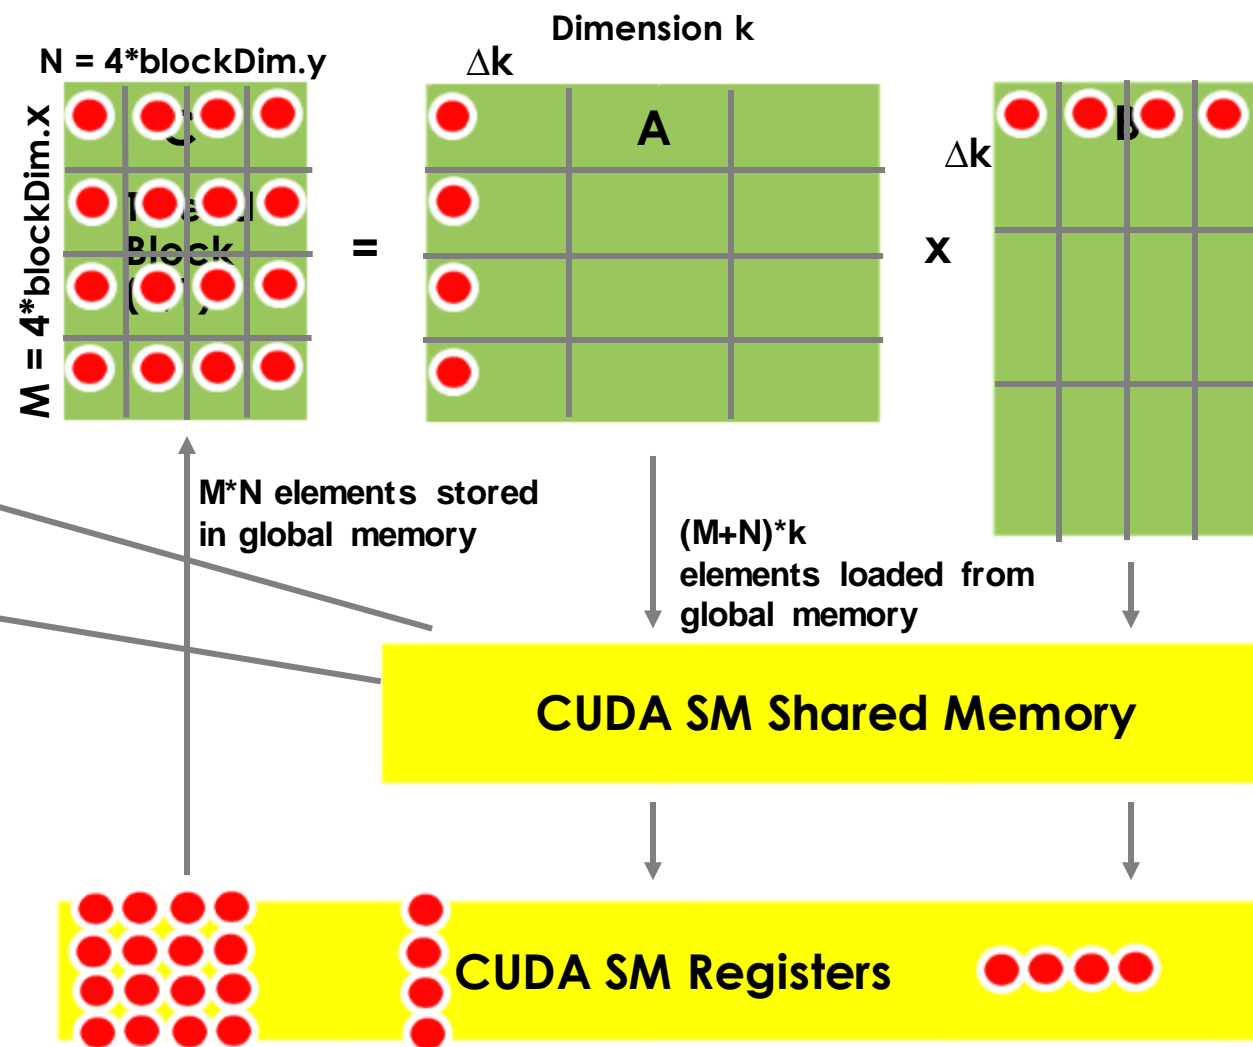
```
    for(int_t m_loc = m_pos + threadIdx.x; m_loc < m_end; m_loc += blockDim.x){
        for(int_t k_loc = k_pos + threadIdx.y; k_loc < k_end; k_loc += blockDim.y){
            lbuf[k_loc - k_pos][m_loc - m_pos] = left[k_loc*m + m_loc];
        }
    }
    }
    Loop: M > blockDim.x;
    Loop: N > blockDim.y;
```

```
    //Load a tile of matrix B(k_pos:TILE_EXT_K, n_pos:TILE_EXT_N):
```

```
    for(int_t n_loc = n_pos + threadIdx.y; n_loc < n_end; n_loc += blockDim.y){
        for(int_t k_loc = k_pos + threadIdx.x; k_loc < k_end; k_loc += blockDim.x){
            rbuf[n_loc - n_pos][k_loc - k_pos] = right[n_loc*k + k_loc];
        }
    }
    }
    __syncthreads();
```

Each CUDA thread block computes:

$C(M = 4 * \text{blockDim.x}, N = 4 * \text{blockDim.y}) += A(M, k) * B(k, N)$



# CUDA BLA Library: +Registers GEMM (algorithm 2)

//Multiply two loaded tiles to produce a tile of matrix

```

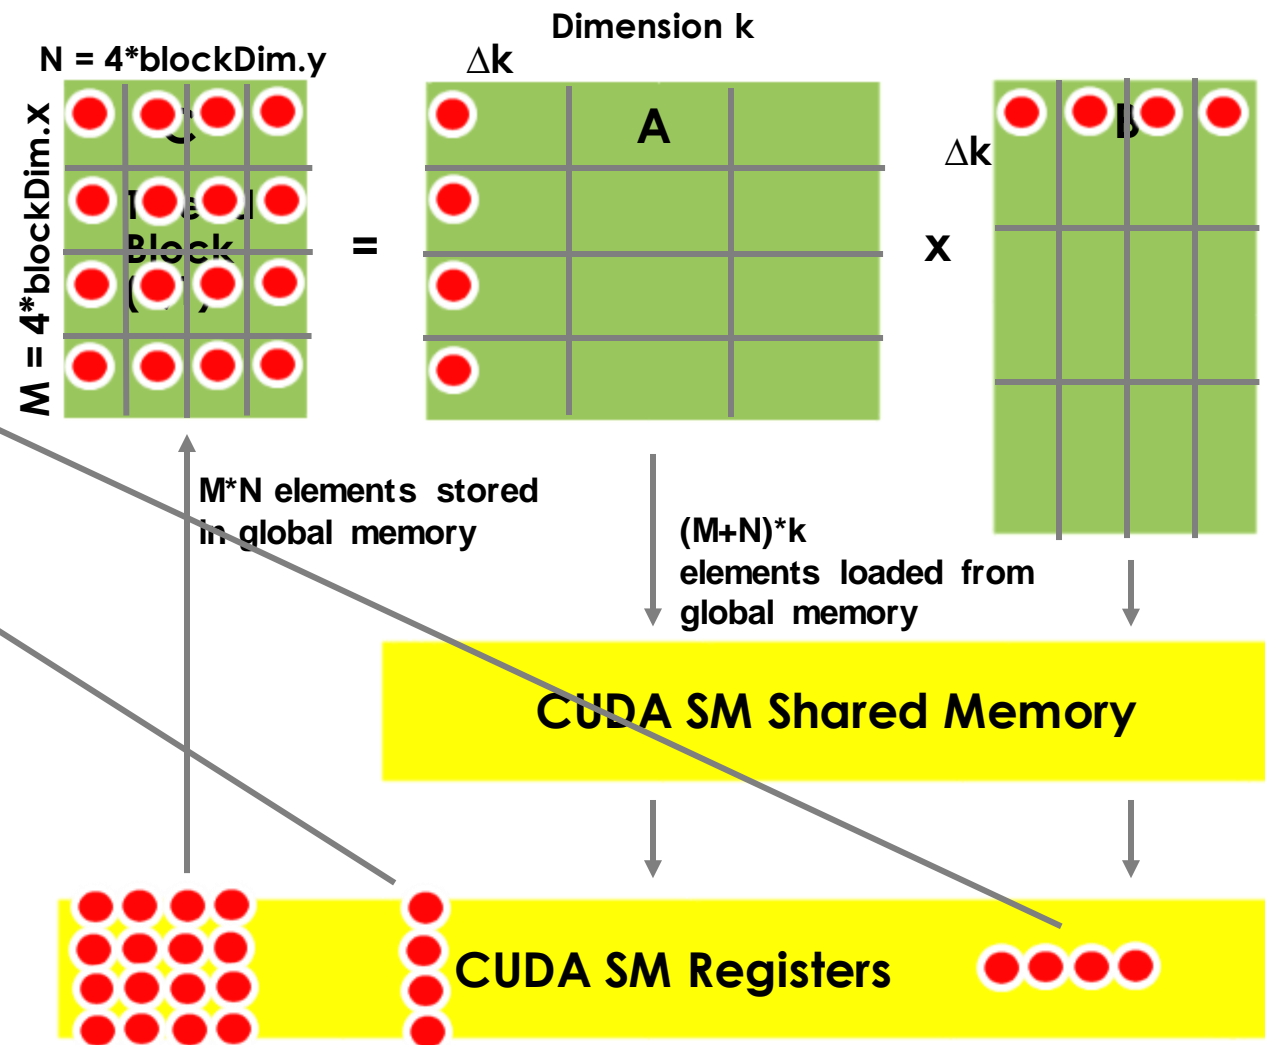
if(k_end - k_pos == TILE_EXT_K){
#pragma unroll
for(int_t l = 0; l < TILE_EXT_K; ++l){
#pragma unroll
for(int_t j = 0; j < 4; ++j) rreg[j] = rbuf[threadldx.y + blockDim.y*j][l];
#pragma unroll
for(int_t j = 0; j < 4; ++j) lreg[j] = lbuf[l][threadldx.x + blockDim.x*j];
#pragma unroll
for(int_t j = 0; j < 4; ++j){
for(int_t i = 0; i < 4; ++i){
dreg[j][i] += lreg[i] * rreg[j];
}
}
}
}
}
else{
for(int_t l = 0; l < (k_end - k_pos); ++l){
#pragma unroll
for(int_t j = 0; j < 4; ++j) rreg[j] = rbuf[threadldx.y + blockDim.y*j][l];
#pragma unroll
for(int_t j = 0; j < 4; ++j) lreg[j] = lbuf[l][threadldx.x + blockDim.x*j];
#pragma unroll
for(int_t j = 0; j < 4; ++j){
for(int_t j = 0; j < 4; ++j){
for(int_t i = 0; i < 4; ++i){
dreg[j][i] += lreg[i] * rreg[j];
}
}
}
}
}
__syncthreads();
    
```

Load registers

Load registers

Each CUDA thread block computes:

$C(M = 4 \cdot \text{blockDim.x}, N = 4 \cdot \text{blockDim.y}) += A(M, k) * B(k, N)$



# CUDA BLA Library: +Registers GEMM (algorithm 2)

//Multiply two loaded tiles to produce a tile of matrix

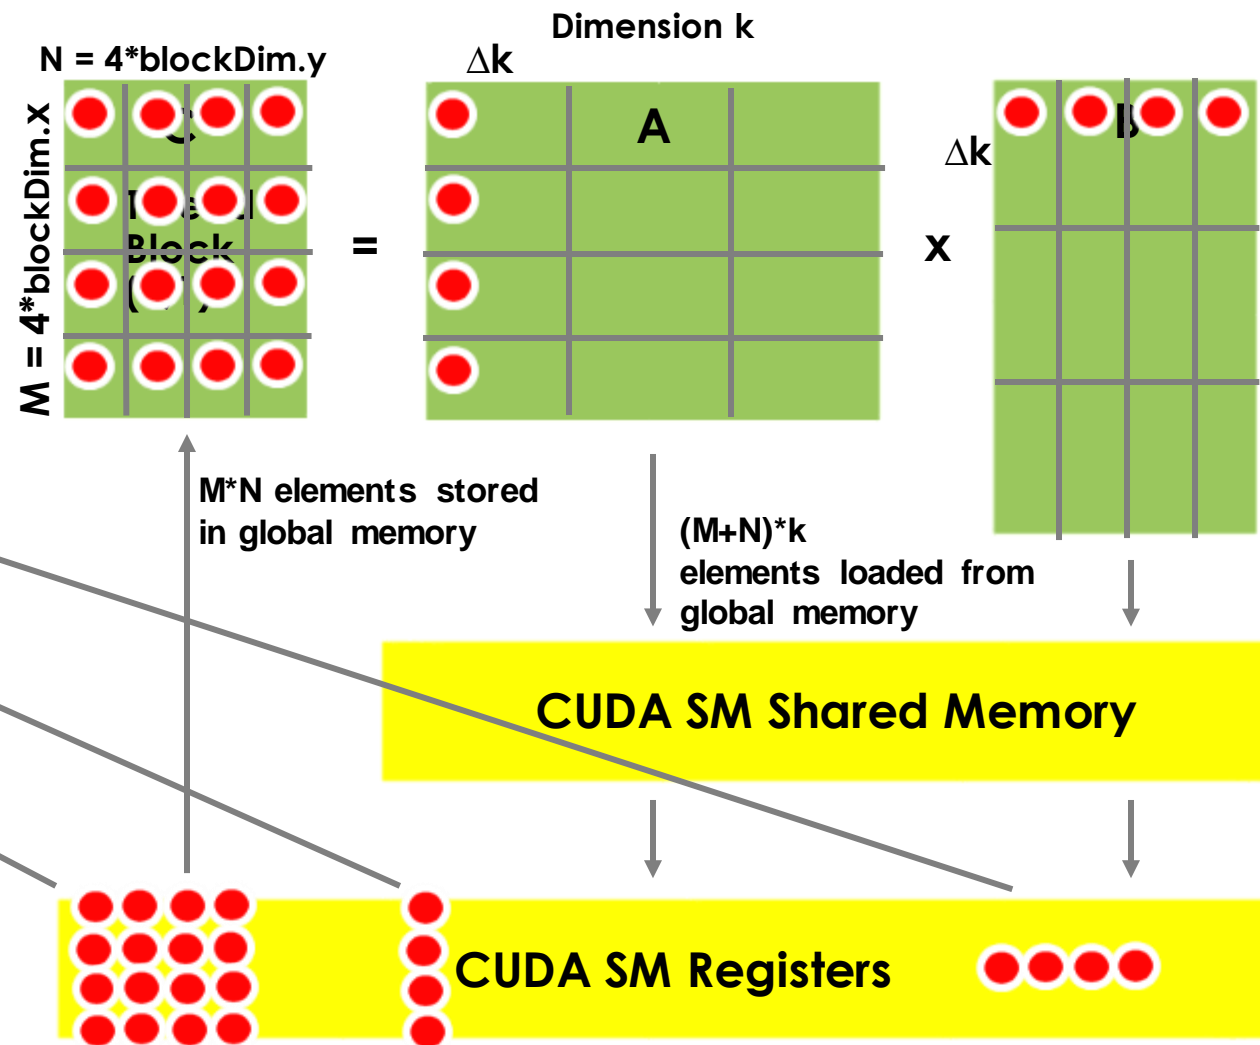
```

if(k_end - k_pos == TILE_EXT_K){
#pragma unroll
for(int_t l = 0; l < TILE_EXT_K; ++l){
#pragma unroll
for(int_t j = 0; j < 4; ++j) rreg[j] = rbuf[threadldx.y + blockDim.y*j][l];
#pragma unroll
for(int_t j = 0; j < 4; ++j) lreg[j] = lbuf[l][threadldx.x + blockDim.x*j];
#pragma unroll
for(int_t j = 0; j < 4; ++j){
#pragma unroll
for(int_t i = 0; i < 4; ++i){
dreg[j][i] += lreg[i] * rreg[j];
}
}
}
}
else{
for(int_t l = 0; l < (k_end - k_pos); ++l){
#pragma unroll
for(int_t j = 0; j < 4; ++j) rreg[j] = rbuf[threadldx.y + blockDim.y*j][l];
#pragma unroll
for(int_t j = 0; j < 4; ++j) lreg[j] = lbuf[l][threadldx.x + blockDim.x*j];
#pragma unroll
for(int_t j = 0; j < 4; ++j){
#pragma unroll
for(int_t i = 0; i < 4; ++i){
dreg[j][i] += lreg[i] * rreg[j];
}
}
}
}
__syncthreads();
    
```

**4x4 matrix outer product from registers by each thread**

Each CUDA thread block computes:

$C(M = 4 * \text{blockDim.x}, N = 4 * \text{blockDim.y}) += A(M, k) * B(k, N)$





# CUDA BLA Library: +Registers GEMM (algorithm 2)

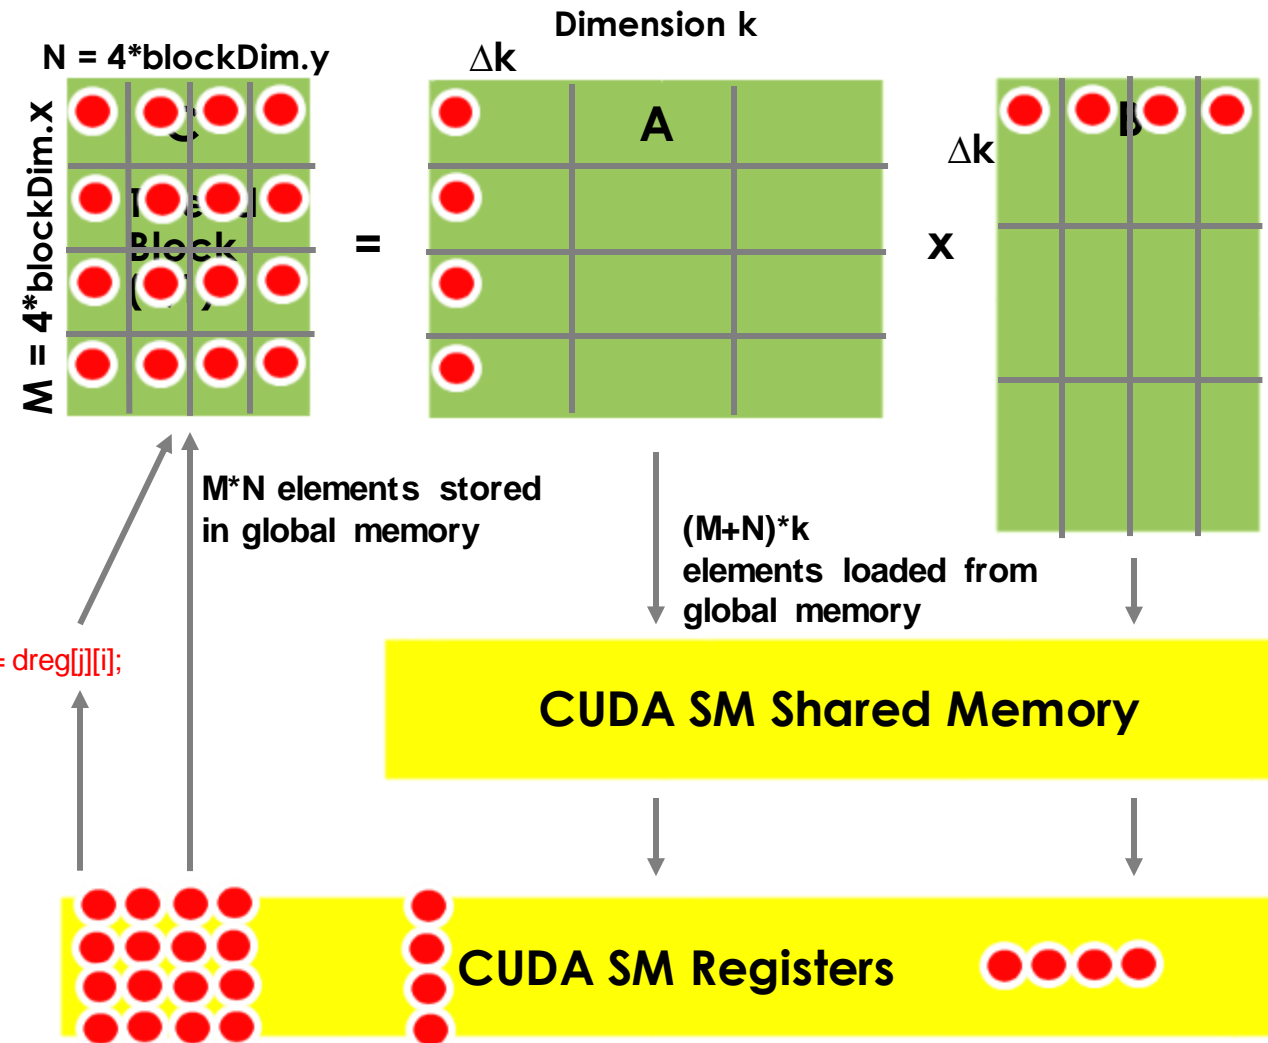
Each CUDA thread block computes:

$C(M = 4 \cdot \text{blockDim.x}, N = 4 \cdot \text{blockDim.y}) += A(M, k) * B(k, N)$

//Store elements of the C matrix in global memory:

```
#pragma unroll
for(int_t j = 0; j < 4; ++j){
#pragma unroll
for(int_t i = 0; i < 4; ++i){
    dest[(n_pos + threadIdx.y + blockDim.y*j)*m + (m_pos + threadIdx.x + blockDim.x*i)] += dreg[j][i];
}
}
```

Upload (4x4) registers to global memory



# CUDA BLA Library: Implement Your GEMM Algorithms

- You will work inside **bla\_lib.cu** source file directly with CUDA GEMM kernels
- Matrix multiplication **{false,false}** case (already implemented):
  - $C(m,n) += A(m,k) * B(k,n)$
  - CUDA kernels: **gpu\_gemm\_nn, gpu\_gemm\_sh\_nn, gpu\_gemm\_sh\_reg\_nn**
- Matrix multiplication **{false,true}** case (your exercise):
  - $C(m,n) += A(m,k) * B(n,k)$
  - CUDA kernels: **gpu\_gemm\_nt, gpu\_gemm\_sh\_nt, gpu\_gemm\_sh\_reg\_nt**
- Matrix multiplication **{true,false}** case (your exercise):
  - $C(m,n) += A(k,m) * B(k,n)$
  - CUDA kernels: **gpu\_gemm\_tn, gpu\_gemm\_sh\_tn, gpu\_gemm\_sh\_reg\_tn**
- Matrix multiplication **{true,true}** case (your exercise):
  - $C(m,n) += A(k,m) * B(n,k)$
  - CUDA kernels: **gpu\_gemm\_tt, gpu\_gemm\_sh\_tt, gpu\_gemm\_sh\_reg\_tt**

# CUDA BLA Library Implementation Benchmark

Testing your BLA GPU kernel implementation (main.cpp: use\_bla() function):

```
for(int repeat = 0; repeat < 2; ++repeat){ //repeat experiment twice
    C.zeroBody(0); //set matrix C body to zero on GPU#0
    bla::reset_gemm_algorithm(0); //choose your algorithm: {0,1,2,7}
    std::cout << "Performing matrix multiplication C+=A*B with BLA GEMM brute-force ... ";
    double tms = bla::time_sys_sec(); //timer start
    C.multiplyAdd(false,false,A,B,0); //default case {false,false}: You goal is {false,true}, {true,false}, {true,true}
    double tmf = bla::time_sys_sec(); //timer stop
    std::cout << "Done: Time = " << tmf-tms << " s: Gflop/s = " << flops/(tmf-tms)/1e9 << std::endl;
    //Check correctness on GPU#0:
    C.add(D,1.0f,0); //adding the correct result with a minus sign (matrix D) should give you zero matrix
    auto norm_diff = C.computeNorm(0); //check its norm
    std::cout << "Norm of the matrix C deviation from correct = " << norm_diff << std::endl;
    if(std::abs(norm_diff) > 1e-7){ //report if norm is not zero enough
        std::cout << "#FATAL: Matrix C is incorrect, fix your GPU kernel implementation!" << std::endl;
        std::exit(1);
    }
}
```

This benchmark is run for all available BLA GEMM algorithms: 0, 1, 2, 7 for the {false,false} case. Your goal is to implement and run other cases: {false,true}, {true,false}, {true,true}! Use TEST\_BLA\_GEMM\_BRUTE, TEST\_BLA\_GEMM\_SHARED, TEST\_BLA\_GEMM\_REGISTER switches in main.cpp:use\_bla() to turn on/off specific GEMM algorithms (0,1,2, respectively).